

ZUSE Automat Agent: Empirical Law Discovery in Elementary Cellular Automata

Abstract

We present ZUSE, a deterministic discovery pipeline for elementary cellular automata (ECA). It runs fixed simulation protocols, evaluates seven operational cycle laws, stores multi-seed WorldRecords, and measures basin fragility without a language model in the discovery loop. Across 20 worlds, ZUSE separates law coverage, observer artifacts, and fragility regimes hidden by coarse visual taxonomy.

The strongest case study concerns periodic-background oscillators in `rule_73/rule_109`. A length-8 sweep finds a $T=15$ mechanism with a five-state locking cycle. A 25-cell, 12-step causal cone reveals an algebraic normal form (ANF) gradient: monomial counts decay with distance from the defect center with $R^2=0.998197$, while active-output degree follows $\text{degree}=24-d+\epsilon$. The gradient generalizes to external length-9/10 $T=15$ witnesses and a background-conditioned `rule_73` holdout. On 219 fragile Q8 targets induced by the long-period replay, 43,425 unit interventions obey an exact topological law: an added state restores vertex or edge redundancy if and only if it bypasses every corresponding critical cut. An exhaustive 404,054-pair audit proves minimum cardinality two in 69/126 vertex and 68/139 edge strata. Exhaustive audits of 3,061,466 triples and 24,362,850 quadruples resolve every remaining collective stratum: the minimum is exactly three in 41/57 vertex and 40/71 edge strata, and exactly four in the remaining 16/16 and 31/31. The underlying global replay covers 5,783,040 configurations and independently confirms 3,296 stationary oscillators omitted by the historical $T \leq 16$ detector.

An exact motif audit then classifies all 27,828,370 candidate pairs, triples, and quadruples into 12 unlabeled Hamming-1 graph motifs and decomposes 1,476 metric-specific minimal rescues by cut coverage and internal-edge dependence. Those verified rescues compile into 265 sparse unit-cost QUBO models with 19,100 variables accumulated across independent models, 32,861 nonzero terms, and exactly 1,476 certified ground states. No quantum hardware is used and no quantum speedup or practical advantage is claimed. Within the 122 rescues from 24 k_2 instances containing both observed mechanisms, the fixed budget of 14 external Hamming-1 incidences is allocated differently: external-attachment rescues have lower mean connectivity toward the period-specific candidate universe and higher mean connectivity toward the disjoint bridge graph than internal-edge-dependent rescues. The bridge-graph difference has the same sign in all 24 instances. This is a conditioned geometric description of the frozen census, not a causal or externally validated classifier.

All scripts, reports, PDFs, and versioned releases are public. Claims remain empirical and protocol-bounded: ZUSE is a reproducible evidence engine for CA law discovery, not a universal autonomous scientist.

1. Introduction

Elementary cellular automata are among the simplest systems known to exhibit complex behavior. A radius-1, binary, one-dimensional CA is fully specified by a single integer from 0 to 255, yet even within this minimal space, Wolfram's empirical taxonomy finds four qualitatively distinct dynamic classes: uniform, periodic, locally chaotic, and complex. Cook's proof that Rule 110 supports universal computation establishes that complexity in ECA is not merely visual: it has computational consequences.

Two questions remain largely open after Wolfram's program. First, the taxonomy is qualitative and coarse: all complex rules fall into Class 4 regardless of their intra-class differences in structure type, periodicity, fragility, or scale behavior. Second, the boundary between the dynamics of the rule and the properties of the

measurement instrument is rarely made explicit. When an observer reports that a run contains gliders, it conflates the CA physics with the heuristic that defined the glider label.

ZUSE Automat Agent addresses both questions through a deterministic, policy-driven discovery loop. The agent runs ECA worlds, applies a fixed stack of heuristic observers, evaluates seven binary cycle laws, and stores multi-seed evidence in persistent world records. No language model participates in the discovery loop: law proposals, world selection, and evidence evaluation are all deterministic. Language-model assistance is restricted to post-run interpretation and documentation.

The result is an empirical atlas of 20 worlds with seven operational categories, measured fragility along two axes (f_{total} and f_{core}), and explicit characterization of two observer artifacts. The atlas is not a new taxonomy of Wolfram's classes; it is a finer-grained, evidence-based map of a 20-world sample that separates cycle-level laws, world-level regimes, and pipeline behavior.

1.1 Contributions

We make the following contributions:

1. **ZUSE Automat Agent** -- A deterministic, policy-driven discovery loop for ECA that accumulates multi-seed law evidence across worlds without symbolic regression or LLM guidance in the loop. The agent combines persistent world-record history, a dedup-gated observer stack, and a seven-law evaluator into a single reproducible pipeline.
2. **A seven-category empirical atlas of 20 worlds** -- We classify 20 worlds spanning ECA rules, Conway's Game of Life patterns, and synthetic controls into seven operational categories (*frontera-rich-estable*, *periodicidad-global*, *oscilador-local*, *multiregimen-productivo*, *multiregimen-escala-dependiente*, *noise-bounded*, *sin-evidencia-multiregimen*) using law coverage, signature diversity, and fragility. The atlas extends Wolfram's four-class taxonomy by capturing intra-class structure, scale-dependent silencing, and negative-control regimes.
3. **A two-dimensional fragility framework** -- We measure f_{total} and f_{core} separately, defining $f_{\text{gap}} = f_{\text{total}} - f_{\text{core}}$ as a quantitative measure of secondary-law churn. Four distinct mechanisms are identified: stable basin (rule_208/209 , $f_{\text{total}} = 0.000$), productive basin switching (rule_137 , $f_{\text{gap}} = 0.318$), noise-boundary fragility (rule_54 , $f_{\text{core}} = 0.677$), and quiescent-background activation (rule_108 , $f_{\text{gap}} = 0.945$).
4. **rule_108 as the unique stationary local-period-2 ECA oscillator** -- Under an exhaustive protocol (128 quiescent ECA rules, 502 non-zero IC words per rule, $\text{span} \leq 32$, $\text{period} \leq 16$), rule_108 is the only ECA rule that produces stationary local period-2 oscillators. The motif $\#.\# \leftrightarrow \###$ follows algebraically from $f(0,1,0) = f(1,0,1) = 1$ and $f(1,1,1) = 0$, and the rule's left-right symmetry ($f(1,c,r) = f(r,c,1)$) explains why the oscillator does not drift.
5. **A measured separation between ECA dynamics and observer artifacts** -- rule_54 single-bit-IC frames are provably translation-invariant (confirmed by frame identity after shift normalization), while observer dedup counts range from 15 to 24 across IC positions. This non-equivariance is characterized as a pipeline property: absolute structure counts depend on IC context, but law signatures remain stable.
6. **An exact cut-coverage law for local basin robustness** -- In 48 frozen length-8 Q8 intervention cubes, 43,425 unit additions to 219 fragile targets restore vertex or edge redundancy if and only if they bypass every corresponding critical cut. Direct cut removal, fresh max-flow, and the geometric predicate agree without exception. A complete 404,054-pair audit then proves minimum rescue cardinality two in 69/126 vertex-connectivity strata and 68/139 edge-connectivity strata. Exhaustive triple and quadruple audits resolve all remaining collective strata at exactly three or four, closing the minimum-cardinality ladder within the frozen protocol.

7. **A verified motif atlas and exact QUBO compilation** -- All 27,828,370 candidate pairs, triples, and quadruples are assigned to one of 12 exact unlabeled Hamming-1 motifs before outcomes are joined. The 1,476 metric-specific minimal rescues compile into 265 sparse unit-cost QUBOs; direct integer evaluation certifies exactly one ground state per rescue and no spurious or missing solution. This is a reusable optimization encoding, not evidence of quantum advantage. A conditioned audit then identifies the motifs with genuine mechanism variation and, for 122 rescues in 24 mixed k_2 instances, decomposes a fixed 14-incidence Hamming budget between the candidate universe, the disjoint bridge graph, and the residual region.

2. Related Work

ZUSE sits at the intersection of three bodies of prior work: empirical ECA taxonomy, automated law discovery, and LLM-based scientific agents. It is related to each tradition but positioned differently from all three: it extends ECA taxonomy inward (intra-class rather than inter-class), it inverts the discovery framing of symbolic regression systems (fixed evaluators rather than generative hypotheses), and it excludes the LLM from the loop rather than centering it.

2.1 ECA taxonomy and complexity

Wolfram's systematic study of elementary cellular automata established the canonical four-class taxonomy: Class 1 (uniform), Class 2 (periodic), Class 3 (chaotic), and Class 4 (complex) [Wolfram2002]. This taxonomy is qualitative and based on visual inspection of space-time diagrams. ZUSE extends it by measuring intra-class structure: two Class-4 rules (`rule_137` and `rule_54`) differ not only in fragility magnitude but in fragility mechanism, a distinction the four-class taxonomy does not capture.

Cook's proof that Rule 110 supports universal computation [Cook2004] established that ECA complexity has computational consequences beyond visual appearance. But computational class is a coarse lens: `rule_110` and `rule_54` are both Class 4, yet ZUSE finds that their fragility mechanisms are qualitatively different -- productive basin switching versus noise-boundary crossing. `rule_110` appears in the ZUSE atlas as a `multiregimen-productivo` world with `f_total = 0.323` and a stable frontera signature. The computational proof characterizes what `rule_110` *can* compute; ZUSE characterizes what it *typically does* under random initialization.

2.2 Automated scientific discovery

AI Feynman [Udrescu2020] demonstrated symbolic regression over physical datasets, recovering known equations from data with interpretable structure. The contrast with ZUSE is deliberate: AI Feynman uses neural networks to propose candidate laws from continuous-variable data, while ZUSE applies fixed binary evaluators to discrete CA dynamics and accumulates evidence without a generative component. ZUSE is not a symbolic regression system; it is a policy-driven measurement pipeline whose outputs are law signatures, not formulas.

More broadly, systems such as Eureqa [Schmidt2009] and recent LLM-based discovery agents frame discovery as hypothesis generation followed by verification. ZUSE inverts this framing: laws are fixed a priori, and the discovery consists of finding which worlds satisfy them and under what conditions. This makes every accepted law signature verifiable from deterministic scripts, at the cost of not proposing new laws automatically.

2.3 ZUSE as evidence engine, not LLM scientist

Recent work on LLM-based scientific agents (e.g., The AI Scientist [Lu2024]) demonstrates that language models can propose hypotheses, design experiments, and write papers with minimal human intervention.

ZUSE occupies a different position in this space: the language model is explicitly excluded from the discovery loop and restricted to post-run interpretation and documentation.

This separation is a design choice, not a limitation. It means that the atlas findings are fully reproducible from the deterministic loop code, and that language-model involvement can be audited at the documentation layer without contaminating the empirical results. The cost is that ZUSE cannot propose new laws; the benefit is that every accepted law has a transparent, non-generative provenance.

2.4 Computational mechanics, particles, and causal states

Computational mechanics provides a formal language for extracting predictive structure from observed sequences. Crutchfield and Young introduced statistical complexity as a way to infer minimal equations of motion from measurement sequences [Crutchfield1989]. Shalizi and Crutchfield later formalized causal states and epsilon-machines as minimal predictive representations of stochastic processes [Shalizi2001]. In cellular automata, this line connects to the broader edge-of-chaos and evolved-computation literature: Langton's lambda experiments framed emergent computation near phase transitions [Langton1990], and Mitchell, Hraber, and Crutchfield studied genetically evolved CAs for global computational tasks [Mitchell1993]. Hanson and Crutchfield applied the domain/particle viewpoint directly to one-dimensional cellular automata by constructing domain filters, particles, and particle interactions for ECA `rule_54` [Hanson1997].

ZUSE uses this vocabulary cautiously. The Fase 68 audit reports a compressed causal-state proxy over the defect trajectory, not a full CSSR reconstruction and not an epsilon-machine. Its negative result should therefore be read as a failure of a deliberately compressed symbolization (`dominant_context`, `defect_size_bucket`), not as a negative result for computational mechanics itself.

2.5 Periodic backgrounds, gliders, and mechanistic audits

The oscillator results in Sections 7.5--7.41 are closest in spirit to the domain/particle tradition in cellular automata: localized structures are studied relative to a background, and the relevant mechanism may be a defect trajectory rather than a raw frame pattern. Lindgren and Nordahl applied complexity measures to one-dimensional cellular automata using block entropy and stochastic finite automata [Lindgren1988]. Boccara, Nasser, and Roger catalogued particle-like structures and interactions in one-dimensional deterministic CA rules [Boccara1991]. Later Rule 54 work developed this line through collision catalogues and logical gates [Martinez2006], and through a formal language for particles on periodic backgrounds [Redeker2010]. Hanson and Crutchfield's `rule_54` analysis demonstrates how regular domains can be filtered to expose particles and interactions [Hanson1997]. ZUSE differs by using exhaustive finite sweeps over prescribed periodic backgrounds and IC words, then auditing the discovered witnesses with Boolean ANF, perturbation, and residual tests.

The ANF analysis uses standard algebraic normal form of Boolean functions: each finite Boolean map is represented as an XOR-sum of monomials over $\text{GF}(2)$. The paper does not claim a new theory of ANF; it uses ANF as an exact finite representation of the 25-input causal-cone maps. The mechanistic terms introduced in the later audit sections, such as ablation, local circuit, and faithful mechanism, are borrowed by analogy from mechanistic interpretability work in neural networks, where small algorithmic circuits are reverse engineered by interventions and decomposition [Elhage2021]. Here those terms refer only to finite ECA truth tables and causal cones.

2.6 Symbolic regression and postponed formula search

Symbolic regression systems such as AI Feynman [Udrescu2020] and PySR [Cranmer2023] search for compact expressions that fit measured data. ZUSE uses symbolic regression only as an external diagnostic. A PySR run is reported as a negative result in Section 9.5, and later symbolic regression is postponed until the dataset contains enough independent comparable measurements to avoid overfitting. This is consistent with

the paper's main design constraint: measurement and acceptance are deterministic, while formula search remains an optional post hoc tool.

3. System: ZUSE Automat Agent

ZUSE Automat Agent is a deterministic discovery loop over cellular automaton worlds. A world is a simulator plus an initial-condition protocol and a time window. For ECA worlds, the simulator is the standard binary radius-1 update rule with periodic boundary conditions. The agent runs a world, computes frame metrics, extracts candidate structures, evaluates a fixed set of laws, updates world-level history, and chooses the next action through a transparent policy.

3.1 Inputs and outputs

A single agent cycle takes as input:

- **World identifier:** the ECA rule number (`0..255`) or a named synthetic world.
- **IC protocol:** either a random seed for standard runs or a designed IC vector for controlled experiments.
- **Width and steps:** fixed integers, typically `width = 64` and `steps = 24..200`, depending on the world's scale protocol.

The outputs of a single cycle are:

- **Frame metrics:** density mean, entropy mean, temporal transition rate, gzip ratio, and mutual information mean.
- **Analysis status:** `ok` or `ruido_no_analizable` after the noise gate.
- **Structure records:** raw `Estructura` outputs with type labels and span information, plus `dedup_structure_count`.
- **Law signature:** frozenset of accepted law names, or the empty set if the run is noise-gated.

3.2 Loop structure

The loop has five layers:

1. **Simulation.** ECA frames are generated from explicit initial conditions with fixed `width`, `steps`, and `seed` or with designed ICs for controlled experiments. The simulator itself is not learned.
2. **Frame metrics.** Each run is summarized by density, entropy, temporal transition rate, gzip compressibility, and temporal mutual information. These features support both individual laws (`complejidad_alta`, `frontera_temporal`, `temporal_scale_stability`) and later meta-analysis.
3. **Observers and deduplication.** A stack of heuristic observers converts frame histories into `Estructura` records with type labels such as `glider`, `bloque`, and `oscilador`. The raw observer outputs are intentionally kept for audit, while `deduplicate_structures` estimates the number of physical structures. The production noise gate uses `dedup_structure_count > 40`.
4. **Cycle-law evaluation.** Seven laws are evaluated on each analyzable run. The result is a law signature: the set of accepted laws for that cycle. Noise-gated runs skip law evaluation rather than forcing a low-confidence signature.
5. **Policy and memory.** The agent stores a persistent `WorldRecord` per world and chooses the next action after each cycle.

3.3 State: WorldRecord

Each world maintains a `WorldRecord` with the following fields relevant to atlas construction:

field	description
<code>visit_count</code>	total cycles run on this world
<code>scores</code>	per-cycle scores used by the policy
<code>noise_count / noise_fraction</code>	count and fraction of noise-gated cycles
<code>law_signatures</code>	list of accepted law signatures as frozensets
<code>unique_law_signature_count</code>	count of distinct non-empty signatures
<code>non_empty_signature_visit_count</code>	visits where at least one law was accepted
<code>law_signature_diversity</code>	unique non-empty signatures divided by non-empty visits, reported after at least five non-empty visits
<code>peak_signature_diversity</code>	maximum clean diversity observed so far
<code>has_multiregime_evidence</code>	monotone boolean, set once peak diversity exceeds 0.5 under low noise
<code>params_tried</code>	tested (<code>steps</code> , <code>width</code> , <code>law_signature</code>) tuples
<code>max_ok_steps / first_noise_steps</code>	scale boundary diagnostics

The important design choice is that empty signatures are retained for audit but excluded from diversity. A world is not multi-regime merely because it alternates between laws and silence; multi-regime evidence requires multiple non-empty law signatures.

3.4 Journal

Every cycle result is appended to a JSONL journal (`outputs/experiments_*/journal_*.jsonl`). Each line is a self-contained JSON record with cycle identifier, world identifier, `steps`, `width`, frame metrics, analysis status, structure counts, law signature, action taken, and the previous `WorldRecord` state visible to the policy at decision time. The journal is the primary reproducibility artifact: the atlas in Section 5, the fragility measurements in Section 6, and the case studies in Section 7 are all derived from journal queries and controlled follow-up scripts.

3.5 Policy

The policy selects the next action from four options:

- **Vary seed** (`repeat_vary_seed`): run the same world with a new IC. Used when the world has a new law signature or confirmed multi-regime evidence and the current cycle is productive.
- **Increase scale** (`increase_steps`): raise `steps` to test scale-dependent behavior. Used when the current world produces analyzable signal and has not reached a known noise boundary.
- **Change world** (`change_world`): move to the next world. Used when the current world is noise-bounded, reaches a known noise boundary, has exhausted repeats, or converges to unproductive silence at maximum scale.
- **Stop by exhaustion**: after the requested cycle budget, persist state and journal artifacts.

The policy has no learned parameters. Its thresholds are fixed constants or explicit guards in code: `dedup_structure_count > 40` for the noise gate, signature diversity `> 0.5` for multi-regime evidence, `noise_fraction < 0.20` for clean diversity, and at least five non-empty visits before diversity is reported.

3.6 Non-generative design

The agent is deliberately non-generative inside the loop. No LLM proposes laws, selects worlds, or evaluates a cycle. Symbolic regression was used only outside the loop for calibration and analysis; it is not part of the online discovery policy. The LLM-assisted work reported here occurs after runs are complete: it helps design

follow-up experiments, interpret artifacts, and write documentation. This separation is important because every accepted law signature in the atlas can be reproduced from deterministic scripts.

4. Seven Cycle Laws

The seven laws are the primary evidence units linking raw ECA frames to the world categories, fragility scores, and observer artifacts reported in Sections 5-8. Each law is evaluated per run -- one world, one initial condition, one step count -- once the observer and dedup pipeline reports `analysis_status = ok`. The output is binary: accepted or rejected. Law signatures are frozensets of accepted law names.

4.1 Design rationale

The seven laws were chosen to span distinct aspects of CA behavior using measurements available from a single fixed-length run. They divide into two groups by input type:

- **Structure-observer laws** (`velocidad_constante`, `periodicidad`, `densidad_estable`, `tipo_unico`): depend on the output of the heuristic observer stack. They can only fire if the run is analyzable (`analysis_status = ok`) and the observers detect at least one structure.
- **Frame-metric laws** (`complejidad_alta`, `frontera_temporal`, `temporal_scale_stability`): depend only on summary statistics computed directly from the frame array, without reference to individual structures.

This split is intentional. Frame-metric laws are symmetry-agnostic: they fire regardless of where structures are in the lattice, how many there are, or whether the observers identify them correctly. Structure-observer laws are richer but carry the observer's heuristic assumptions. Section 8 reports two cases where structure-observer laws produce artifacts that frame-metric laws do not.

All laws produce binary output (accepted / rejected). The choice of binary rather than continuous output is also intentional: it makes signatures comparable across runs and worlds without requiring score normalization, and it forces explicit calibration of each threshold.

4.2 Formal criteria

#	Law	Inputs	Criterion	Constants
1	<code>velocidad_constante</code>	Position tracks of moving structures	At least 50% of moving tracks (<code>velocity > 0.05</code> cells/step) have linear $x(t)$ with normalized residual < 0.15	-
2	<code>periodicidad</code>	Structure type list	At least one structure classified as <code>oscilador</code>	-
3	<code>densidad_estable</code>	Frame density time series	Coefficient of variation $CV = \sigma(\rho) / \mu(\rho) < 0.15$	-
4	<code>tipo_unico</code>	Structure type set	Exactly one structure type present	-
5	<code>complejidad_alta</code>	Frame metrics	<code>entropy_mean > 0.80</code> and <code>transition_rate > 0.25</code>	-
6	<code>frontera_temporal</code>	Frame metrics	<code>entropy_mean > 0.80</code> and $0.28 < \text{transition_rate} < 0.4352$	upper threshold calibrated 2026-05-24
7	<code>temporal_scale_stability</code>	Frame metrics + steps	<code>temporal_load = steps * gzip_ratio / transition_rate < 19.03</code>	threshold calibrated 2026-05-24

`temporal_scale_stability` rejects any run with `transition_rate = 0` (quiescent or static configurations), since temporal load is undefined (`infinity`).

4.3 Calibrated constants

Neither threshold can be derived analytically: the boundary between organized frontier dynamics and pure chaos has no closed form in ECA. Both constants were set empirically on real ECA runs and are valid within the atlas protocol (`width = 64`, `steps` roughly `24..200`).

The `frontera_temporal` upper threshold `0.4352` is the midpoint between the maximum `transition_rate` observed for `rule_110` (`0.4147`) and the minimum for `rule_30` (`0.4557`) across six canonical seeds at `steps = 24`, `width = 64`.

The `temporal_scale_stability` threshold `19.03` was fit on `datasets/fase2c_v3.csv` (120 ECA scale samples). A decision tree at `max_depth = 4` achieved accuracy `0.908`, precision `0.886`, and recall `0.954` on the `analysis_ok` label.

4.4 Caveats

`tipo_unico` is an observer-dependent exploratory signal, not a mirror-invariant physical property. Fase 6b showed that `rule_110` and `rule_124` are left-right mirrors of each other with identical dynamics, yet `tipo_unico` can fire asymmetrically depending on orientation. `tipo_unico` is retained in the atlas for its exploratory value but should not be used as evidence of physical asymmetry.

`frontera_temporal` and `temporal_scale_stability` both depend on `transition_rate`. Fase 4a and later tree analyses identify transition rate as the main discriminator separating organized frontier dynamics from pure chaos or static order. Other metrics (`density_mean`, `gzip_ratio`, `mutual_info_mean`) are useful context features but should not be treated as independent causal evidence without ablation.

These caveats do not weaken the atlas: they clarify which signals reflect physical ECA dynamics and which reflect the current observer design. Section 8 returns to both artifacts with controlled experiments.

4.5 Law signatures and the atlas

A law signature is a frozenset of accepted law names for one run. The empty frozenset is valid and indicates a run that passed the noise gate but accepted no law. Law signatures are the unit of evidence in the atlas: the world categories in Section 5 are defined by how signatures distribute across seeds and scales, and the fragility measurements in Section 6 count how often signatures change under perturbation.

`frontera_temporal` is a proper subset of `complejidad_alta` by construction (it adds the upper bound on transition rate). Any run that accepts `frontera_temporal` also accepts `complejidad_alta`; the converse is not required. This containment is visible in the law coverage matrix: every `yes` in the `frontera_temporal` column co-occurs with a `yes` in the `complejidad_alta` column.

5. World Atlas: 20 Worlds and Dynamic Categories

The atlas is derived from `outputs/world_taxonomy/law_map.md`. It contains 20 worlds: ECA rules, designed synthetic controls, and Life-like controls. Each world is summarized by law coverage, non-empty visit ratio, noise ratio, signature diversity, mean law count, dominant signature, and measured fragility where available.

The atlas is not a score table. A high `mean_laws` value, high signature diversity, and low fragility mean different things. The taxonomy therefore separates five positive dynamic families from two bookkeeping categories: `noise-bounded` for worlds stopped by the dedup gate, and `sin-evidencia-multiregimen` for controls or worlds without sufficient evidence for one of the positive families.

5.1 Category definitions

category	operational signal	representative worlds
frontera-rich-estable	low signature diversity, high stable law richness (mean_laws ≥ 4.0)	rule_46, rule_208, rule_209
periodicidad-global	global period-2 behavior; periodicidad in nearly all non-empty visits	rule_51
oscilador-local	bounded local period-2 structure on a quiescent background	rule_108
multiregimen-productivo	multiple non-empty law signatures with productive visits	rule_18, rule_54, rule_109, rule_110, rule_124, rule_137
multiregimen-escala-dependiente	real signature diversity but most high-scale visits become analyzable silence	rule_90
noise-bounded	pre-law failure under the deduplicated structure gate	rule_30, rule_150
sin-evidencia-multiregimen	no sufficient evidence of multi-regime or stable-rich behavior in the current protocol	life_blinker, life_block, life_glider, synthetic_bloque, synthetic_glider, synthetic_oscilador

world	category	mean_laws	peak_diversity	f_total	f_core
rule_208	frontera-rich-estable	6.000	0.167	0.000	0.000
rule_209	frontera-rich-estable	6.000	0.167	0.000	0.000
rule_46	frontera-rich-estable	5.833	0.333	0.031	0.031
rule_51	periodicidad-global	4.500	0.333	0.193	0.000
rule_108	oscilador-local	2.000	0.167	0.992	0.047
rule_90	multiregimen-escala-dependiente	0.500	0.600	0.172	0.000
rule_110	multiregimen-productivo	2.727	0.600	0.323	0.198
rule_124	multiregimen-productivo	2.167	0.600	0.224	0.083
rule_109	multiregimen-productivo	2.000	0.667	0.307	0.307
rule_18	multiregimen-productivo	2.308	0.800	0.349	0.135
rule_137	multiregimen-productivo	2.867	0.833	0.630	0.312
rule_54	multiregimen-productivo	1.917	0.800	0.714	0.677
rule_30	noise-bounded	1.100	0.000	0.021	0.021
rule_150	noise-bounded	0.750	0.000	0.023	0.023
life_blinker	sin-evidencia-multiregimen	3.000	0.200	1.000	1.000
life_block	sin-evidencia-multiregimen	2.000	0.200	0.016	0.016
life_glider	sin-evidencia-multiregimen	2.357	0.333	0.032	0.032
synthetic_bloque	sin-evidencia-multiregimen	2.000	0.200	n/a	n/a
synthetic_glider	sin-evidencia-multiregimen	3.167	0.400	n/a	n/a
synthetic_oscilador	sin-evidencia-multiregimen	2.286	0.400	n/a	n/a

This classification is intentionally operational. A world can be reclassified if a wider protocol produces different evidence; the atlas records what the current deterministic protocol has measured.

The same principle applies to `frontera_temporal` candidates. Fase 20a found many additional ECA rules that are rich in `frontera_temporal` at the fixed sweep scale (`steps = 24`), but Fase 20b showed that the strongest four new candidates do not remain `frontera-rich-estable` under long-journal policy scaling. Fase 20c therefore treats `frontera-short-scale` as a candidate tier, not as an atlas-grade category. The atlas promotes worlds only when short-scale richness survives the broader validation protocol.

5.2 Law coverage

The law coverage matrix uses seven columns:

```

velocidad_constante
periodicidad
densidad_estable
tipo_unico
complejidad_alta
frontera_temporal
temporal_scale_stability

```

Each cell has one of four states: accepted in the dominant signature or in at least half of non-empty visits (`yes`), observed but below half (`partial`), never observed in non-empty visits (`-`), or unknown because no non-empty visits exist (`?`).

Cell states:

- `yes`: law appears in the dominant signature or in at least 50% of non-empty visits.
- `partial`: law appears in at least one non-empty visit but in less than 50%.
- `-`: non-empty visits exist and the law never appears.
- `?`: no non-empty visits.

world	vel	per	den	tipo	compl	front	tss
<code>life_blinker</code>	-	yes	yes	yes	-	-	-
<code>life_block</code>	-	-	yes	yes	-	-	-
<code>life_glider</code>	partial	-	yes	yes	-	-	-
<code>rule_108</code>	-	yes	-	yes	-	-	-
<code>rule_109</code>	-	-	yes	partial	yes	yes	yes
<code>rule_110</code>	yes	-	yes	-	yes	yes	partial
<code>rule_124</code>	partial	-	yes	-	yes	yes	yes
<code>rule_137</code>	partial	-	yes	partial	yes	yes	yes
<code>rule_150</code>	-	-	yes	-	yes	-	yes
<code>rule_18</code>	yes	-	-	yes	yes	-	yes
<code>rule_208</code>	yes	-	yes	yes	yes	yes	yes
<code>rule_209</code>	yes	-	yes	yes	yes	yes	yes
<code>rule_30</code>	-	-	yes	-	yes	-	yes
<code>rule_46</code>	yes	-	yes	yes	yes	yes	yes
<code>rule_51</code>	-	yes	yes	yes	yes	-	yes
<code>rule_54</code>	yes	-	partial	partial	yes	-	yes
<code>rule_90</code>	partial	-	partial	-	partial	-	yes

world	vel	per	den	tipo	compl	front	tss
synthetic_bloque	-	-	yes	yes	-	-	-
synthetic_glider	yes	-	yes	yes	-	-	partial
synthetic_oscilador	-	yes	-	yes	-	-	partial

Law coverage matrix for the 20-world atlas

Dark = accepted; light = rejected; mid = partial/observer-dependent marker.



Figure 1. Law coverage matrix for the 20-world atlas. Dark cells indicate laws accepted in the dominant signature or in at least half of non-empty visits; mid cells indicate partial activation; light cells indicate rejection.

The matrix reveals three broad patterns:

- Synthetic and Life-like controls validate observer semantics.** `life_blinker` and `synthetic_oscilador` activate `periodicidad`; block-like worlds activate `densidad_estable` and `tipo_unico`; synthetic gliders activate `velocidad_constante`.

2. **Class-4 and frontier worlds separate into distinct families.** `rule_137`, `rule_110`, `rule_124`, `rule_109`, and `rule_54` are multi-regime worlds with two or three dominant laws. By contrast, `rule_46`, `rule_208`, and `rule_209` activate six of seven laws with low diversity.
3. **periodicidad is IC-family sensitive.** Under random ICs it appears in designed controls, in the global complement rule (`rule_51`), and in the local oscillator (`rule_108`), but not in the complex frontier worlds. Under explicitly periodic ICs, however, Fase 21a finds production `periodicidad` in 207/256 ECA rules. The law is therefore not dead or ECA-inaccessible; it is controlled by the IC family. This is why Section 7 treats `rule_108` separately rather than folding it into ordinary stable-rich behavior.

5.3 Key atlas rows

The following rows anchor the category structure:

world	category	mean_laws	peak_diversity	dominant signature
<code>rule_208</code>	<code>frontera-rich-estable</code>	6.000	0.167	<code>velocidad_constante + densidad_estable + tipo_unico + complejidad_alta + frontera_temporal + temporal_scale_stability</code>
<code>rule_209</code>	<code>frontera-rich-estable</code>	6.000	0.167	<code>velocidad_constante + densidad_estable + tipo_unico + complejidad_alta + frontera_temporal + temporal_scale_stability</code>
<code>rule_46</code>	<code>frontera-rich-estable</code>	5.833	0.333	<code>velocidad_constante + densidad_estable + tipo_unico + complejidad_alta + frontera_temporal + temporal_scale_stability</code>
<code>rule_137</code>	<code>multiregimen-productivo</code>	2.867	0.833	<code>densidad_estable + complejidad_alta + frontera_temporal</code>
<code>rule_54</code>	<code>multiregimen-productivo</code>	1.917	0.800	<code>complejidad_alta + temporal_scale_stability</code>
<code>rule_51</code>	<code>periodicidad-global</code>	4.500	0.333	<code>periodicidad + densidad_estable + tipo_unico + complejidad_alta + temporal_scale_stability</code>
<code>rule_108</code>	<code>oscilador-local</code>	2.000	0.167	<code>periodicidad + tipo_unico</code>
<code>rule_90</code>	<code>multiregimen-escala-dependiente</code>	0.500	0.600	<code>temporal_scale_stability</code>

These rows show why the taxonomy cannot be reduced to a single richness score. `rule_208` and `rule_209` are maximally rich and stable; `rule_137` is less rich but highly diverse; `rule_108` is law-sparse but category-defining because it is the only local oscillator; `rule_90` has high diversity evidence but low non-empty yield because its high-scale visits become silent.

5.4 Scientific role of the atlas

The atlas is the bridge between cycle-level laws and world-level claims. A law signature describes one run. A world category describes how signatures behave across seeds, scales, and perturbations. This distinction is what makes later fragility measurements interpretable: $f_{total} = 0.630$ in `rule_137` means something different from $f_{total} = 0.992$ in `rule_108` because the atlas identifies different category-defining cores.

6. Fragility: `f_total`, `f_core`, `f_gap`

6.1 Protocol

Fragility is measured by exhaustive one-bit IC perturbation. For each measured world and canonical seed, every bit in the IC is flipped individually, and the resulting run is evaluated through the full pipeline (simulation, frame metrics, observers, dedup, law evaluation). The reference is the law signature of the unperturbed run.

Protocol parameters:

- **IC width:** 64 for most fragility measurements; 128 for the designed `rule_108` local-oscillator IC.
- **Perturbations per seed:** one per bit position (64 or 128, depending on IC width).
- **Seeds per world:** usually 3 canonical seeds, giving 192 perturbations for width-64 worlds. Designed-IC worlds such as `rule_108` use a canonical IC rather than random seeds.
- **Steps:** world-specific canonical steps (e.g., 24 for `rule_46`, 48 for `rule_137`, 96 for `rule_54`).

6.2 Metrics

Three primary fragility metrics are defined:

- `f_total`: fraction of perturbations that produce a different law signature from the reference (including noise-gated runs and silence).
- `f_core`: fraction that changes the category-defining core laws. Noise and silence count as core changes because the defining regime is lost.
- `f_gap = f_total - f_core`: secondary-law churn. Perturbations that change the signature without affecting the core-defining laws.

A fourth component is tracked separately:

- `f_noise`: fraction of perturbations that produce `analysis_status = ruido_no_analizable` (noise-gate crossing).

`f_noise` is a component of `f_total` and `f_core`; it is reported separately because it identifies a specific observer-boundary mechanism.

6.3 Core-law convention

Core laws are defined per category:

category	core laws
<code>frontera-rich-estable</code>	the full six-law frontier signature
<code>periodicidad-global</code>	<code>periodicidad</code>
<code>oscilador-local</code>	<code>periodicidad</code> and <code>tipo_unico</code>
<code>multiregimen-productivo</code>	the reference signature of that specific seed
<code>multiregimen-escala-dependiente</code>	<code>temporal_scale_stability</code>

For `multiregimen-productivo` worlds, the reference signature varies by seed. `f_core` is therefore computed per seed against that seed's reference and then averaged.

6.4 Fragility spectrum

Fase 22 completes the physical fragility spectrum for all measured non-synthetic worlds in the atlas. Synthetic controls are excluded from `f_total` and `f_core` because they are frame generators rather than evolved

systems with a perturbable initial condition. The completed spectrum includes ECA worlds, Life fixtures, stable-rich frontier worlds, noise-bounded productive pockets, and the designed `rule_108` local-oscillator IC.

world	category	f_total	f_core	f_gap	f_noise	mechanism
<code>rule_208</code>	frontera-rich-estable	0.000	0.000	0.000	0.000	stable basin
<code>rule_209</code>	frontera-rich-estable	0.000	0.000	0.000	0.000	stable basin
<code>life_block</code>	sin-evidencia-multiregimen	0.016	0.016	0.000	0.000	stable Life fixture
<code>rule_30</code>	noise-bounded	0.021	0.021	0.000	0.000	productive pocket
<code>rule_150</code>	noise-bounded	0.023	0.023	0.000	0.000	productive pocket
<code>rule_46</code>	frontera-rich-estable	0.031	0.031	0.000	0.000	stable basin
<code>life_glider</code>	sin-evidencia-multiregimen	0.032	0.032	0.000	0.000	stable Life fixture
<code>rule_90</code>	multiregimen-escala-dependiente	0.172	0.000	0.172	0.000	secondary churn
<code>rule_51</code>	periodicidad-global	0.193	0.000	0.193	0.000	secondary churn
<code>rule_124</code>	multiregimen-productivo	0.224	0.083	0.141	0.000	productive switching
<code>rule_109</code>	multiregimen-productivo	0.307	0.307	0.000	0.000	productive switching
<code>rule_110</code>	multiregimen-productivo	0.323	0.198	0.125	0.000	productive switching
<code>rule_18</code>	multiregimen-productivo	0.349	0.135	0.214	0.000	productive switching
<code>rule_137</code>	multiregimen-productivo	0.630	0.312	0.318	0.000	productive switching
<code>rule_54</code>	multiregimen-productivo	0.714	0.677	0.037	0.375	noise-boundary
<code>rule_108</code>	oscilador-local	0.992	0.047	0.945	0.000	quiescent-background activation
<code>life_blinker</code>	sin-evidencia-multiregimen	1.000	1.000	0.000	0.000	periodic fixture disruption

The spectrum is category-aligned at the extremes: `frontera-rich-estable` occupies the low end, while `multiregimen-productivo` occupies the upper ECA range. The `life_blinker` control reaches `f_total = 1.000` because any single-cell perturbation breaks the exact Life oscillator fixture. `rule_108` remains the main ECA structural outlier: `f_total = 0.992` with only `f_core = 0.047`.

6.5 Fragility mechanisms

The atlas identifies several distinct mechanisms by which one-bit IC perturbations change law signatures:

Stable basin (`rule_208`, `rule_209`): `f_total = 0.000`. All perturbations preserve the reference signature. The basin for the six-law frontier signature is wide enough that no measured single-bit perturbation escapes it. `rule_46` is nearly identical (`f_total = 0.031`).

Stable Life fixture (`life_block`, `life_glider`): perturbations rarely change the reference signature (`f_total ≤ 0.032`) because the fixture remains structurally recognizable after most one-cell flips. This is fixture-level robustness, not evidence of a broad ECA basin.

Productive pocket (`rule_30`, `rule_150`): the worlds are noise-bounded in the long journal, but the non-empty pockets that survive the gate are stable under one-bit perturbation (`f_total ≈ 0.02`). Their category

is defined by frequent pre-law noise at scale, not by fragility of the productive signatures.

Productive basin switching (`rule_137`, and the non-noise-boundary `multiregimen-productivo` worlds): perturbations move the IC among productive law-signature regimes. The world never falls into silence or noise; `f_noise` = 0.000 throughout. `rule_137` is the strongest clean case (`f_total` = 0.630), with more than 80% of perturbations switching regime in the two most fragile measured seeds.

Noise-boundary fragility (`rule_54`): perturbations cross the observer noise gate rather than moving between productive regimes. The mechanism requires complex ICs near the dedup threshold; single-bit ICs from bare backgrounds do not approach the gate (Section 8). `f_noise` = 0.375 makes `rule_54` the only measured world where noise-gate crossings dominate fragility.

Periodic fixture disruption (`life_blinker`): any one-cell perturbation breaks the exact Life period-2 reference signature (`f_total` = 1.000). This is not basin switching or noise-boundary crossing; it is the brittleness of a minimal periodic fixture under full-grid perturbation.

Quiescent-background activation (`rule_108`): the canonical IC has only two active bits on a zero background. Nearly any background perturbation ignites new dynamics and changes secondary laws, producing `f_total` = 0.992. The core oscillator survives unless the perturbation lands near the motif (`f_core` = 0.047). The result is the largest `f_gap` in the atlas (0.945): nearly all fragility is secondary, not core.

6.6 The `f_core` / `f_gap` separation


The main result of the fragility analysis is the separation between core and secondary law changes:

- `rule_51` (`periodicidad-global`): `f_total` = 0.193, `f_core` = 0.000. Global periodicity survives all measured perturbations; only secondary laws (`densidad_estable`) toggle.
- `rule_108` (`oscilador-local`): `f_total` = 0.992, `f_core` = 0.047. The local oscillator survives nearly all perturbations; secondary laws are maximally sensitive.
- `rule_54` (`multiregimen-productivo`): `f_total` = 0.714, `f_core` = 0.677. The core productive signature changes frequently; `f_gap` = 0.037 is near zero.
- `rule_137` (`multiregimen-productivo`): `f_total` = 0.630, `f_core` = 0.312. Both core and secondary transitions are common; `f_gap` is approximately equal to `f_core`.

These four cases span the space of possible (`f_core`, `f_gap`) combinations. Together they demonstrate that `f_total` alone is insufficient: two worlds can have similar total fragility with opposite core/secondary decompositions.

Fragility spectrum

One-bit IC fragility separates total disruption from core-law disruption.



Labels show representative extremes; synthetic controls are omitted from the scatter.

Figure 2. Fragility spectrum. The scatter separates total one-bit perturbation sensitivity (f_{total}) from disruption of the category-defining law core (f_{core}).

7. Case Studies


7.1 rule_108 -- Unique Local Oscillator

Discovery and formal profile

rule_108 was identified during a targeted local-oscillator search (Fase 16) using minimal ICs on a quiescent background ($f(0,0,0) = 0$). The canonical IC is a pair of active cells separated by one gap ($\#.\#$, word 101 in binary). Under rule_108, this IC produces an exact period-2 local oscillator (Figure 3): the gap fills in each step ($\#.\# \rightarrow \###$) and then empties again, repeating indefinitely with zero drift.

rule_108 stationary local oscillator

Seed `##` alternates with `###` on a quiescent background ($T=2$).



Orange box marks the active local support used by the stationary oscillator baseline.

Figure 3. *rule_108 stationary local oscillator*. Active cells are dark and quiescent cells are light; the orange box marks the bounded local support.

The oscillator is stationary (center of mass fixed), bounded ($\text{span} \leq 3$), and stable over 200 steps with zero drift on a uniform-zero background.

Formal profile (6 canonical seed labels, `width = 128`, `steps = 200`, `IC = pair_gap1`): `periodicidad` and `tipo_unico` are accepted in 6/6 runs. Mean `dedup_structure_count` = 1.000. The oscillator is deterministic given the canonical IC; the seed labels are retained only to keep the profile format consistent with other atlas worlds.

Algebraic derivation

The oscillator follows from three entries in the `rule_108` table:

Neighborhood	Output	Role
010	1	isolated active center stays active
101	1	gap fills in: <code>##</code> -> <code>###</code>
111	0	center empties: <code>###</code> -> <code>##</code>

`rule_108` is left-right symmetric ($f(l, c, r) = f(r, c, l)$), which explains why the oscillator does not drift: the two flanking cells exert equal influence on the center.

Fragility: quiescent-background activation

`rule_108` has the largest fragility gap in the atlas:

Metric	Value
<code>f_total</code>	0.992
<code>f_core</code>	0.047
<code>f_gap</code>	0.945
Core positions	61, 62, 63, 65, 66, 67
Pattern	clustered

The mechanism is geometric. The canonical IC (pair_gap1, 2 active bits in width = 128) leaves more than 120 cells at zero. A one-bit perturbation at any of those background positions activates the quiescent background: because $f(0,1,0) = 1$, an isolated 1 on a zero background immediately grows, producing new detectable structures. This changes secondary laws while leaving periodicidad and tipo_unico intact as long as the oscillator core is undisturbed. A perturbation within the core neighborhood (positions 61..63, 65..67) displaces or destroys the oscillator, accounting for $f_{core} = 0.047$.

This constitutes a third fragility mechanism, distinct from productive basin switching (rule_137, Section 7.3) and noise-boundary fragility (rule_54, Section 7.2): quiescent-background activation. The core behavior is robust, but the minimal IC makes secondary laws highly sensitive to background activation.

Uniqueness

Fase 18 ran an exhaustive search over all 128 ECA rules with quiescent backgrounds ($f(0,0,0) = 0$), testing 502 non-zero IC words per rule (all non-empty binary words of length 1..8), with width = 128, steps = 200, burn-in of 80 steps, and requiring zero drift (stationary oscillators only). Only rule_108 produced stationary local period-2 oscillators; no other rule produced a stationary local oscillator for periods 2..16 and span <= 32.

A companion sweep with the stationarity requirement relaxed found a distinct family of moving local oscillators -- see Section 7.5.

The longer-IC extension strengthens the same conclusion. Extending the stationary sweep to all 7,676 non-zero IC words of length 9..12 (982,528 rule/IC runs) produced 3,802 detections, all in rule_108. No new stationary oscillator rule appears beyond the length-1..8 baseline; the minimum witness is still the embedded 101 motif (000000101 at length 9).

The family is internal to rule_108: 132 out of 179 candidate IC words are accepted by the production observer as periodicidad, with oscillator spans 3, 5, 6, 7, and 8. All confirmed oscillators have period exactly 2; no longer period was found.

7.2 rule_54 -- Noise Gate Anatomy

rule_54 is the clearest example of noise-boundary fragility: perturbations do not merely move the run to another productive signature, but can push the observer output across the deduplicated structure gate. The production gate is:

```
dedup_structure_count > 40 -> ruido_no_analizabile
```

Fase 13: anatomy of the gate

Fase 13 measured three productive rule_54 ICs at steps = 96 and perturbed each by all 64 one-bit flips. The reference deduplicated counts were close to the gate:

seed	reference dedup	noisy flips / 64
20260638	32	14
20260640	33	18
20260642	39	40

Across the three seeds, 72/192 flips crossed into ruido_no_analizabile ($f_{noise} = 0.375$). Every noisy flip crossed for the same reason: dedup_structure_count > 40. No alternative noise mechanism was observed.

The sensitive positions formed a clustered, multi-hot pattern rather than a single contiguous block. Bins near the periodic boundary (0..7 and 56..63) were repeatedly implicated, and bit 5 was the only bit whose flip

crossed the gate in all three measured seeds.

Fase 19: controlled single-bit negative case

Fase 19 tested whether bit 5 was a special absolute coordinate of `rule_54`. It replaced the complex ICs with controlled single-bit ICs: for each `k = 0..63`, the initial state contained only one active bit at position `k`.

The result separates CA physics from the observer pipeline:

- The ECA frames are translation-invariant after shift normalization.
- The observer/dedup counts are not translation-equivariant for this wide-spreading pattern:
`dedup_structure_count` ranges from 15 to 24 across positions.
- The law signature is identical for all 64 positions: `temporal_scale_stability`.
- Every single-bit IC remains far below the gate (`dedup <= 24 < 40`).

Therefore, bit 5 is not a privileged coordinate of `rule_54`. The Fase 13 signal arises from the interaction between complex IC geometry and the observer/gate pipeline. Complex ICs close to the threshold can be pushed across it by local flips; a single active cell cannot.

Mechanism

`rule_54` has high total and core fragility (`f_total = 0.714`, `f_core = 0.677`), but its mechanism differs from `rule_137`. In `rule_137`, perturbations tend to move between productive regimes. In `rule_54`, a large fraction of perturbations cross an analysis boundary: the run becomes too fragmented for the current observer/dedup gate.

This makes `rule_54` a methodological case study as much as a dynamical one. It shows that the atlas can identify worlds whose measured fragility is dominated by proximity to an observer threshold. It also motivates the caveat that absolute structure counts should not be treated as symmetry-invariant physical observables without equivariance checks.

7.3 rule_137 -- Productive Basin Switching

`rule_137` is the primary example of productive basin switching: one-bit IC perturbations change the law signature without ever crossing the noise gate or reaching silence. All fragility is productive (`f_noise = 0.000`, `f_silence = 0.000`), making it the cleanest case in the atlas for inter-basin transitions.

Fragility profile

Three canonical seeds at `steps = 48`, `width = 64`:

seed	reference signature	f_total
20260633	<code>complejidad_alta + frontera_temporal</code>	0.812
20260635	<code>complejidad_alta + densidad_estable + frontera_temporal</code>	0.219
20260673	<code>complejidad_alta + densidad_estable + frontera_temporal + temporal_scale_stability + tipo_unico + velocidad_constante</code>	0.859

Aggregate: `f_total = 0.630`, `f_core = 0.312`, `f_gap = 0.318`.

The per-seed range (`0.219..0.859`) is the widest in the atlas. Even the least fragile measured seed has `f_total > 0.2`. The most fragile seeds (20260633 and 20260673) flip on more than 80% of one-bit perturbations.

Mechanism

`f_noise = 0.000`: no perturbed IC crosses the deduplicated structure gate. The world remains analyzable throughout. The fragility is a property of productive basin geography, not proximity to an observer threshold.

The pattern is `dispersed`: sensitive positions are distributed across the IC width rather than concentrated near a motif. This is consistent with a world that has many narrow productive basins whose boundaries intersect throughout the IC space.

`peak_diversity = 0.833` -- the highest in the atlas. The canonical seeds themselves already visit multiple distinct productive regimes. The fragility measurement extends this: not just that the world can reach different signatures under different seeds, but that a single-bit perturbation to any one canonical IC is enough to move between regimes.

`f_core` and `f_gap` interpretation

`f_core = 0.312` reflects genuine regime switching: flips that remove or change laws defining the canonical signature. `f_gap = 0.318` reflects secondary-law churn: the core productive signature survives, but laws on the signature boundary (such as `densidad_estable` or `tipo_unico`) toggle.

The two components are roughly equal (0.312 vs 0.318), meaning `rule_137` sits in a region where both core-regime transitions and secondary-law transitions are common. This is structurally different from `rule_108` (`f_gap = 0.945`, where the core oscillator is robust and secondary laws dominate) and from `rule_54` (`f_gap = 0.037`, where the core productive signature changes but almost no fragility is secondary).

Contrast with `rule_54` and `rule_108`

`rule_54` and `rule_137` both have high `f_total` (0.714 vs 0.630), but the mechanisms are opposite: `rule_54` fragility is dominated by noise-gate crossings (`f_noise = 0.375`), while `rule_137` fragility is entirely productive. A perturbed `rule_137` IC stays analyzable and law-rich; it is simply in a different productive regime.

`rule_108` contrasts from the opposite direction: its `f_core = 0.047` shows that the defining behavior (the local oscillator) is nearly indestructible, while `rule_137`'s `f_core = 0.312` shows that its defining signatures change under nearly a third of all one-bit flips.

7.4 `rule_46`, `rule_208`, `rule_209` -- Stable-Rich Frontier

These three worlds define the `frontera-rich-estable` category: low signature diversity, near-maximal law richness, and very low fragility. They are the counterexample that revised the early atlas interpretation of `frontera_temporal`.

The first 15-world atlas made `frontera_temporal` look rare: it appeared only as a minority law in class-4 multi-regime worlds such as `rule_137`, `rule_110`, and `rule_54`. Fase 11 showed that this was a sampling artifact. A sweep over all 256 ECA rules (`seeds = 20260523..20260525`, `W = 64`, `T = 24`) found 38 rules where `frontera_temporal` activates in at least two of three seeds, and 17 rules where it activates in all three.

The top rules by law richness were `rule_46`, `rule_208`, and `rule_209`. Formal six-seed profiles (`20260523..20260528`, `W = 64`, `T = 24`) placed all three in a new category:

world	mean laws	peak diversity	category
<code>rule_46</code>	5.833	0.333	<code>frontera-rich-estable</code>
<code>rule_208</code>	6.000	0.167	<code>frontera-rich-estable</code>
<code>rule_209</code>	6.000	0.167	<code>frontera-rich-estable</code>

The dominant signature is the same six-law set:

```
velocidad_constante
densidad_estable
tipo_unico
complejidad_alta
frontera_temporal
temporal_scale_stability
```

Only `periodicidad` is absent. This makes the family nearly maximal under the current seven-law system without relying on multi-regime exploration.

Stable richness rather than multi-regime diversity

The category is defined by the conjunction of high richness and low diversity. `rule_137` is rich because it moves among several productive law signatures. The frontier-rich worlds are rich because the same high-law signature appears reliably across seeds.

This distinction matters operationally. A policy that only looks for signature diversity would miss these worlds, even though they produce more accepted laws per visit than any multi-regime world in the atlas. The Fase 11 taxonomy update therefore adds:

```
frontera-rich-estable := mean_laws >= 4.0 and peak_diversity <= 0.5
```

evaluated after noise-bounded and multi-regime cases.

Complement symmetry and independent convergence

`rule_46` and `rule_209` are a complement pair ($46 = 255 - 209$): exchanging zeros and ones maps one into the other. Their shared profile is therefore one physical phenomenon seen through global bit inversion.

`rule_208` is more surprising. Its complement is `rule_47`, not `rule_46` or `rule_209`, yet it reaches the same maximum-richness profile. This suggests that the `frontera-rich-estable` regime is not a single isolated symmetry orbit; at least two distinct ECA regions converge to the same six-law frontier.

Fragility

Fase 12 measured one-bit fragility for the three worlds using the same protocol as `rule_137` and `rule_54`:

world	f_total	f_core	f_noise
rule_46	0.031	0.031	0.000
rule_208	0.000	0.000	0.000
rule_209	0.000	0.000	0.000

The result is the opposite end of the fragility spectrum from `rule_137`. Where `rule_137` has many narrow productive basins ($f_{total} = 0.630$), `rule_208` and `rule_209` have measured basins so wide that no single-bit flip changes the law signature. `rule_46` is only slightly fragile: two of 192 single-bit perturbations change signature across the three measured seeds.

This confirms that high law richness does not imply high fragility. Richness can arise either from many neighboring productive regimes (`rule_137`) or from a single broad, stable regime (`rule_46/208/209`).

Scientific revision

The correct conclusion is not that `frontera_temporal` is intrinsically rare. It is rare in the original discovery atlas because the original world sequence under-sampled stable high-richness boundary worlds. In the full

ECA sweep, `frontera_temporal` is a robust marker of the `frontera-rich-estable` family.

7.5 Moving Oscillator Family -- Minimal Period-2 Gliders

A companion sweep to Fase 18 searched all 128 quiescent ECA rules ($f(0,0,0) = 0$) for oscillators that translate at constant velocity. The detection protocol matched Fase 18 except: the zero-drift requirement was replaced by a requirement of constant non-zero drift confirmed over three consecutive periods; `width` was extended to 256 and `steps` to 300 to give moving patterns room to travel. The 502 non-zero IC words of length 1..8 were tested per rule (64,256 total runs, 312 s).

Eight rules produce moving oscillators: `rule_6`, `rule_20`, `rule_38`, `rule_52`, `rule_134`, `rule_148`, `rule_166`, `rule_180`. All share the same minimal glider pattern:

Step	Active shape (normalized offsets)
t	[0] -- single active cell
t+1	[0, 1] -- two adjacent active cells
t+2	[0] displaced by +/-2


The oscillator alternates between one and two active cells while traveling at speed 1 cell per step -- the maximum velocity for a radius-1 ECA. Mean active span per period is 0.5 (alternating span 0 and span 1). All eight rules are confirmed `edge_touch = False` within `width = 256`.

Figure 5. Moving oscillator (glider) -- period T=2, speed 1.

rule_20 (drift +2)	rule_6 (drift -2)	
t=0: . . . 1	 1 . . .	
t=1: . . . 1 1	 1 1 . . .	
t=2: 1 . . .	. . . 1	<- period boundary
t=3: 1 1 . .	. . 1 1	
t=4: 1 .	. 1	<- period boundary

Active cells shown as 1, quiescent background as . . Each period advances the pattern two positions in the travel direction.

Figure 6. Moving oscillator (glider) -- period $T=2$, speed 1



Left: rule_20 (right-moving, drift +2 per period). Right: rule_6 (left-moving, drift -2 per period). Active cells dark; quiescent background light. Rows are time steps $t=0..7$. The oscillator alternates $[0] \leftrightarrow [0,1]$ while translating one cell per step.

Structure: four mirror pairs

The eight rules form four left-right symmetric pairs:

Left-moving	Right-moving	b5	b7
rule_6	rule_20	0	0
rule_38	rule_52	1	0
rule_134	rule_148	0	1
rule_166	rule_180	1	1

Bits **b5** (neighborhood 101) and **b7** (neighborhood 111) vary across the family but do not participate in the glider cycle: neither neighborhood occurs during quiescent travel with a two-cell active pattern. The eight rules are therefore a single physical family parametrized by two inactive bit choices and the left-right direction.

Contrast with rule_108

Property	rule_108	Moving family
Drift per period	0	± 2
Period T	2	2
Mean active span	~ 2	0.5
Speed	0 (stationary)	1 (maximum)
Rules in scope	1	8 (4 mirror pairs)

The two families do not overlap. rule_108 does not appear among the eight moving rules, and none of the eight moving rules produce stationary local oscillators under Fase 18. Together, the two sweeps partition the quiescent local oscillator landscape into a unique stationary oscillator (rule_108) and a unique minimal glider family (8 rules, 4 mirror pairs).

Uniqueness within the protocol

No IC word of length $1..8$ produced a moving oscillator with $T > 2$ or $|\text{drift}| \neq 2$ in any quiescent rule. The minimal glider $[0] \leftrightarrow [0, 1]$ at period 2 and maximum speed is the only moving local oscillator family under this protocol.

The longer-IC extension reaches the same result. Testing all 7,676 non-zero IC words of length 9..12 over all 128 quiescent rules (982,528 rule/IC runs) produced 2,059 moving detections, all within the same eight-rule family. No new moving-oscillator rule appears. The minimal witnesses are still the old glider seeds embedded in longer words (000000001 for right movers, 000000010 for left movers). The sweep also filtered 9,822 period-1 moving particle aliases across 32 rules; these are moving particles, not internal period-2 oscillators. Extensions to IC words longer than 12 remain open (Section 10.2). The non-zero background extension is reported in Section 7.6.

The shift from quiescent backgrounds to periodic backgrounds is the point at which the atlas system becomes a mechanism-discovery instrument rather than a catalogue alone. Sections 7.1--7.5 show what the fixed ZUSE protocol can separate under zero-background assumptions: stationary oscillators, moving period-2 gliders, and observer artifacts. Section 7.6 asks the same atlas question under a richer background model. That single protocol change exposes oscillators that are invisible in the quiescent regime, including the rule_73/rule_109 family that later drives the T=15 and ANF-gradient audits. Thus the long ANF block below is not a separate paper appended to the atlas: it is a high-resolution audit of the strongest mechanism surfaced by the atlas sweep itself.

7.6 Periodic-background oscillator sweep

The two sweeps above restrict the background to quiescent zero cells. A third sweep tests whether replacing the background with a non-zero periodic tiling changes the oscillator landscape.

Protocol. All 256 ECA rules are tested against 15 unique non-zero periodic backgrounds with template lengths 1, 2, and 4. Each rule/background pair is tested against 502 non-zero IC words of length 1..8 centered in a width-256 grid for 300 steps with 80 burn-in. The detector identifies exact recurrence of the localized difference between the perturbed run and the unperturbed background orbit; global background periodicity alone is not counted as a local oscillator. Total: 1,927,680 rule/background/IC runs, 122,253 candidate detections.

Result. Thirty rules produce stationary local oscillators under at least one non-zero periodic background; 36 rules produce moving oscillators. Of the 30 stationary rules, 29 are new relative to the zero-background baseline. Of the 36 moving rules, 28 are new.

The following phenomena appear under non-zero backgrounds but not under the quiescent zero background:

- **Period-4 stationary oscillators.** rule_54 and rule_147 produce stationary period-4 oscillators under 0001 background.
- **Period-4 moving oscillator.** rule_180 produces a T=4 glider with drift +4 under 0001 background (shapes [0] -> [0,1] -> [0,2,3] -> [0]), distinct from its T=2 speed-1 glider under quiescent background.
- **Speed-0.5 gliders.** Multiple rules (including rule_3, rule_17, rule_27, rule_35, rule_39) produce T=2 gliders with drift +/-1 under non-zero backgrounds, corresponding to 0.5 cells per step. Under quiescent zero background the only observed glider speed was 1 cell/step.
- **rule_108 under all-one background.** rule_108 appears in the stationary list with background 1 and the same motif ### / #.# as the quiescent result, confirming that the rule_108 oscillator is intrinsic to the rule table.

Separation of regimes. The zero-background uniqueness claims (Sections 7.1 and 7.5) are not contradicted: they describe the quiescent regime. Under zero background, the quiescent local oscillator space is sparse (1 stationary rule, 8 moving rules); under non-zero periodic backgrounds, the landscape expands substantially (30 stationary, 36 moving) and includes period and speed classes absent from the quiescent regime. The two regimes should not be merged.

7.7 Period-8 background oscillator sweep (Fase 24)

A fourth oscillator sweep extended the periodic-background protocol to length-8 primitive binary backgrounds, testing three questions: do longer background periods introduce new oscillator rules, new period classes $T > 4$, or glider speeds outside $\{0, 0.5, 1\}$?

Protocol. All 256 ECA rules are tested against 30 primitive binary necklaces of length 8. A length-8 binary string is primitive if its minimal period is exactly 8; the necklace representative is the lexicographically smallest member of its rotation class. The count of 30 follows from Mobius inversion: $(2^8 - 2^4) / 8 = 30$. Each rule/background pair is tested against 502 non-zero IC words of length 1..8 in a width-256 grid, 300 steps, 80 burn-in. The differential detector and period search window (2..16, span ≤ 32) are unchanged from Section 7.6. Total: 3,855,360 runs.

Result. The sweep produces 323,872 candidate detections after filtering 95,121 period-1 aliases. New stationary rules beyond the length-1/2/4 baseline are `rule_62`, `rule_118`, `rule_131`, and `rule_145` (4 rules, all $T=3$). New moving rules are `rule_7`, `rule_9`, `rule_21`, `rule_25`, `rule_31`, `rule_45`, `rule_61`, `rule_65`, `rule_67`, `rule_75`, `rule_87`, `rule_88`, `rule_89`, `rule_101`, `rule_103`, `rule_111`, `rule_125`, `rule_173`, and `rule_229` (19 rules).

New period classes. Periods $T=6, 8, 10, 12$, and 15 appear for the first time; under backgrounds of length 1, 2, and 4 the maximum observed period was $T=4$. $T=15$ is a non-trivial period not divisible by the background length (8). Fase 24 reports it as an emergent period; Sections 7.8 and 7.9 subsequently identify its rule family and establish its five-state mechanism under F^3 .

New glider speed. Speed $2/3$ cell/step (drift ± 2 , $T=3$) is observed for the first time; the prior speed set was $\{0, 0.5, 1\}$. Representative cases are `rule_9` (drift -2, $T=3$, background `00001001`) and `rule_65` (drift +2, $T=3$, background `00000001`). Two further rules with the same speed signature, `rule_111` and `rule_125`, appear in the candidate table. Direct rule-table reflection, $g(1, c, r) = f(r, c, 1)$, confirms two exact left-right mirror pairs: `rule_9` $\leftrightarrow$ `rule_65` and `rule_111` $\leftrightarrow$ `rule_125`. The paired rules carry opposite drift signs with the same $T=3$ speed magnitude, giving the speed- $2/3$ family the same left/right mirror structure as the speed-1 family in Section 7.5.

Phase dependence. A rotation sub-test applied all 8 rotations of the canonical background to 10 sampled rules while holding the IC fixed. None of the 10 samples is active in all eight rotations after circular-geometry correction. `rule_62` and `rule_118` activate in 7 of 8 background rotations; moving cases range from 6/8 (`rule_9`, `rule_65`) to 2/8 (`rule_111`, `rule_45`). The IC therefore requires a particular alignment with the background to nucleate the oscillator.

Co-translation test (Fase 25). A strict test co-translates both background and IC through $k=0..7$ for the same 10 cases. The background and XOR perturbation orbits are exact translations in 80/80 runs. The original `linear_shape` preprocessing recovers only 58/80 signatures: 22 moving runs cross positions 255/0, causing a false linear span and rejection. Circular shape canonicalization (largest-gap cut plus continuous position unwrapping) recovers 80/80 signatures and all 10/10 cases. Thus the physics and recurrence detector are co-translation equivariant once cyclic geometry is represented correctly. The remaining fixed-IC phase dependence is physical, although the linear observer overstated its severity for moving rules.

Summary. All three Fase-24 questions are answered affirmatively: period-8 backgrounds introduce 4 new stationary rules and 19 new moving rules, expand the period set to include $T=6, 8, 10, 12$, and 15 , and introduce speed $2/3$ cell/step as a new rational class. The zero-background uniqueness claims of Sections 7.1 and 7.5 are unaffected.

7.8 Anatomy of the $T=15$ family (Fase 26)

The longest period observed in Fase 24 was analyzed separately to determine whether it represented a single accidental witness or a coherent family. All 221 $T=15$ detections are stationary and occur in only two rules: `rule_73` (123 detections) and `rule_109` (98). Both rules are left-right symmetric, and black/white conjugation maps each rule exactly to the other. The detections cover 14 primitive length-8 backgrounds, 20 rule/background pairs, and 25 temporal motifs up to cycle phase. The minimum witness is `rule_109` on background `00011001` with the two-cell IC word `01`.

Temporal locking. Every participating unperturbed background enters a temporal orbit of period $T_{bg}=3$ (with transient length 0..2). The localized perturbation has fundamental period $T_{local}=15$, giving the same locking ratio $T_{local}/T_{bg}=5$ in all 20 rule/background pairs. Thus $T=15$ is not inherited directly from the spatial background length 8. Long-horizon reruns of one minimal witness per pair preserve exact recurrence through step 900 in 20/20 cases; the detector scans upward from period 1 and therefore excludes smaller fundamental periods.

Basin width. Persistence in time does not imply robustness to initialization. Holding the IC fixed while rotating the background preserves $T=15$ in only 23/160 runs. One-bit mutations of the 20 minimal IC witnesses preserve it in 4/134 runs. Most failed perturbations settle into shorter localized periods $T=3$ or $T=6$; a small number produce $T=12$ or no localized period in the search window. The family is therefore temporally exact but basin-narrow.

These results establish a persistent background-locked family rather than a single numerical coincidence. Fase 26 measures but does not explain the five-to-one locking ratio; Fase 27 addresses its finite-state mechanism next.

7.9 Five-state locking mechanism (Fase 27)

Fase 26 measured $T_{local}/T_{bg}=5$ across all 20 minimal $T=15$ representatives but left the origin of that ratio open. Fase 27 closes the computational half of the question by tracking the localized XOR defect $D(t) = X(t) \text{ XOR } B(t)$ once per background period.

Protocol. After burn-in, the defect is sampled at $t = 81 + k \cdot T_{bg}$ for $k=0..20$, covering four complete candidate five-cycles for each of the 20 minimal (rule, background, IC) witnesses from Fase 26. Acceptance requires eight simultaneous checks: the background returns to the exact same phase at every sample; the first five defect states are mutually distinct; the fifth transition closes the cycle; no shorter cycle under F^3 explains the sequence; four consecutive cycles repeat in both canonical and raw-position encodings; transitions remain deterministic across cycles; and the oscillator is stationary over every local period.

Results. All eight checks are satisfied for 20/20 representatives. Every $T=15$ oscillator in the family is stationary ($drift=0$). The defect cycles through exactly five distinct states under F^3 , the three-step evolution operator, and five applications of F^3 are necessary and sufficient to restore the complete background-plus-defect configuration. The same result holds over four consecutive cycles in both relative-shape and absolute-position encodings.

Interpretation. The locking ratio $T_{local}/T_{bg}=5$ is the cycle length of the defect under F^3 , not a resonance with the spatial background length 8. The result is a computational state-cycle derivation: it establishes the finite-state mechanism without reducing the five nodes to a closed-form symbolic identity over the rule-table algebra of `rule_73/rule_109`. That symbolic derivation remains open.

7.10 Induced defect rule and exact conjugation (Fase 28)

Fase 28 examines the local algebra inside the five-state cycle. A localized XOR defect does not evolve under the original ECA rule f in isolation. If b is a three-bit background neighborhood and d its XOR-defect neighborhood, the exact induced update is

$$\text{delta}_f(b,d) = f(b \text{ XOR } d) \text{ XOR } f(b).$$

The analysis profiles this induced rule over all 300 microsteps contained in the 100 F^3 edges from the 20 minimal representatives.

Analytical conjugation. Let C denote bitwise complementation. Since rule_{109} is the black/white conjugate of rule_{73} , their global maps satisfy $F_{109}(C(X)) = C(F_{73}(X))$. If both the full state and background are complemented, induction gives $X_{109}(t)=C(X_{73}(t))$ and $B_{109}(t)=C(B_{73}(t))$ for every t . Therefore:

$$D_{109}(t) = C(X_{73}(t)) \text{ XOR } C(B_{73}(t)) = D_{73}(t).$$

At local level, the equivalent identity is $\text{delta}_{109}(C(b),d) = \text{delta}_{73}(b,d)$. Thus the defect orbit is invariant, not complemented, under the simultaneous black/white conjugation. This statement is analytical; exhaustive checks over all 64 local (b,d) combinations and 10/10 complemented orbit pairs serve as implementation validation.

Sparse-support hypothesis rejected. Every one of the 100 F^3 macro-transitions uses all eight ordinary truth-table entries in its causal defect cone. No single induced (b,d) key appears in every macro-transition. The five-cycle therefore cannot be explained by one fixed sparse subset of local entries.

The negative result still exposes phase structure. The sizes of the induced-key intersections for transitions $S_0 \rightarrow S_1$ through $S_4 \rightarrow S_0$ are $[9,6,7,5,4]$ for rule_{73} and $[11,11,10,9,8]$ for rule_{109} . These signatures are rule-specific rather than universal. A closed-form derivation must therefore encode spatial phase or a higher-order block state, rather than truth-table entry presence alone.

7.11 Block-locality limits and background-indexed shape families (Fase 29-30)

Section 7.10 concludes that a closed-form derivation of the five-cycle must encode spatial phase rather than truth-table entry presence alone. Fase 29 and Fase 30 test two complementary refinements of that conclusion.

Defect-shape locality (Fase 29). For each of the 20 minimal representatives and each of the five cycle phases, the canonical XOR-defect shape is computed at sample times $t = 81, 84, 87, 90, 93$. If the $T=15$ cycle were a pure defect-only dynamic, every background would produce the same canonical shape in each phase. Instead, rule_{73} produces 8-9 distinct shapes per phase across its 10 backgrounds, and rule_{109} produces 8. Extending the comparison window outward to fixed local blocks (active defect span plus up to three padding cells on each side) does not help: no nontrivial block signature is shared across all backgrounds in any phase. The $w=0$ result in the block-signature scan shows that the active defect span itself already carries background-dependent context -- no additional padding is needed before the discrimination occurs. Only the trivial token $d000 \rightarrow 0$ (background cells evolving to background under all three microsteps) is universal.

Shape families (Fase 30). Although the defect shapes are background-dependent, they are not unstructured. Treating two five-state cycles as equivalent when one is a cyclic phase rotation of the other, the 20 representatives decompose into 7 distinct families for rule_{73} and 8 for rule_{109} , yielding 13 global families. The largest family has 3 members. Two families are shared across the conjugate rules: one of size 3 (two rule_{73} backgrounds and one rule_{109} background) and one of size 2 (background word 00110101 producing phase-rotationally equivalent defect cycles under both rules). This second coincidence is not implied by the analytical conjugation of Fase 28 -- the bitwise complement of 00110101 is 11001010 , a distinct

word -- so it is an independent structural coincidence. Scalar background descriptors (`active_count`, `transition_count`, `active_transition_pair`) do not determine the family. The canonical temporal orbit of the background under the same rule is exact in this representative set, but this is largely a restatement of full background identity rather than a compact symbolic law.

Interpretation. Fase 29 and Fase 30 together bracket the derivation problem: a correct symbolic account must use the full temporal background orbit (or an equivalent compact encoding), and maps that orbit to one of a finite set of defect-cycle shapes plus a phase offset. A derivation predicting a single universal five-state cycle is falsified by Fase 29; a derivation predicting arbitrary unclustered shape variation is refuted by Fase 30.

7.12 Compact state variable for the T=15 family (Fase 31-32)

Fase 30 reduced the T=15 family to 13 shape families but required the full temporal background orbit as the discriminating descriptor. Fase 31 and Fase 32 search for a shorter description.

Descriptor search (Fase 31). A compact global background descriptor does not exist among the candidates tested (length-2.4 circular subpattern counts, parity, run lengths, and orbit prefixes up to 24 bits). Globally, only `orbit_prefix_24` determines the family, which is substantially a restatement of full background identity. Conditioned on the ECA rule, the shortest non-orbit candidate is the circular multiset of length-4 background subwords (`subpatterns_len4`): it separates all 10 backgrounds per rule into distinct buckets with zero ambiguity. A global decision tree over all features achieves only 0.700 training accuracy at depth 4, confirming that no shallow feature combination separates all 13 families simultaneously.

Rotation generalization (Fase 32). The circular subpattern multiset is invariant under background rotation. Fase 32 tests whether the predicted family is preserved across all seven non-trivial rotations of each of the 20 representative backgrounds, under two modes: `fixed_ic` (background rotated, IC position fixed) and `cotranslated_ic` (background rotated, IC shifted by the same amount to preserve local alignment).

Under `fixed_ic`, only 3 of 140 rotations produce T=15 at all, and 1 matches the predicted family. Under `cotranslated_ic`, all 140 of 140 rotations produce T=15 and all 140 match the predicted family. The co-translation result is not evaluated on new backgrounds but on rotational variants of the known 20 representatives; it validates that the descriptor is rotation-equivariant within the confirmed family set, not that it predicts previously unseen data.

Compact state variable. The complete minimal description consistent with all observations is the triple (`rule`, `subpatterns_len4`, `IC/background alignment`). The rule identity determines which of the two conjugate families applies; the subpattern multiset identifies the shape-family class within that rule; and the IC/background alignment selects the phase within the five-state cycle. A derivation operating on the background word alone without IC alignment is falsified by the `fixed_ic` mode (3/140 detections). Any derivation that correctly maps this triple to a defect cycle shape and phase offset is a complete symbolic account of the T=15 family.

7.13 External length-9/10 validation of the T=15 mechanism (Fase 33-34)

Fase 33 first audits whether the compact descriptor from Section 7.12 can be falsified inside the length-8 background universe. Across all binary circular length-8 backgrounds there are only two collisions of the `subpatterns_len4` descriptor: `00110111/00111011` and `00010011/00011001`. Both collisions are already inside the confirmed T=15 set and both preserve the same defect-cycle family under the corresponding rule. There is no unseen length-8 background, outside rotations of the 20 known representatives, that shares a T=15 `subpatterns_len4` descriptor under the same rule. Thus the natural external test of the descriptor is impossible inside length 8.

Fase 34 therefore moves outside the length-8 universe, but only after a preflight checks whether the prerequisite for the five-to-one locking mechanism still exists. Primitive length-9 and length-10 backgrounds do contain temporal period-three cases under both `rule_73` and `rule_109`: 11 backgrounds per rule at length 9 and 22 per rule at length 10. A targeted validation then tests only these 66 backgrounds, the two rules, and the same 502 non-zero IC words of length 1..8, for 33,132 rule/background/IC runs.

The targeted sweep finds 90 $T=15$ detections across 8 external backgrounds: one length-9 background under `rule_73`, five length-10 backgrounds under `rule_73`, and two length-10 backgrounds under `rule_109`. No length-9 background under `rule_109` produces a $T=15$ witness in this test. These backgrounds are not rotations of the length-8 representatives. The result therefore shows that the $T=15$ mechanism is not an artifact of primitive length-8 backgrounds: it generalizes when $T_{bg}=3$ is preserved. The compact descriptor from Section 7.12 remains a length-8 family identifier; extending that descriptor to variable background length remains a separate symbolic problem.

7.14 Transition tables and effective orbit embeddings (Fase 35-38)

Sections 7.10--7.13 reduce the $T=15$ family to finite defect-cycle shapes, compact length-8 descriptors, and external $T_{bg}=3$ witnesses. Fases 35--38 then move from visual shape families to the explicit macro-operator acting on the localized defect.

Transition tables (Fase 35). For each of the 20 minimal representatives, the five transitions $D_i \rightarrow D_{i+1}$ under F^3 are expanded into explicit local transition tables over the induced defect rule $\delta f(b,d)=f(b \oplus d) \oplus f(b)$. The table signature is never shared across different visual families, so it is a sufficient discriminator. It is not identical to the visual family partition, however: the verdict is `TABLE_REFINES_FAMILY`. Some multi-member families share the same table exactly or up to cyclic phase rotation, while others split by rule or by mechanism. This shows that the fundamental object is the induced macro-transition table, with the visual family as a coarser quotient.

Effective orbit identity (Fase 36). The clearest table identity appears in family `F00` under `rule_109`: the backgrounds `00001001`, `00010011`, and `00011001` have the same exact five-phase transition table. Fase 36 explains this algebraically at the orbit level. Although the three initial backgrounds have different preperiods, by the sampling time they have converged to the same canonical period-3 background orbit $\{00001001, 00101101, 00111111\}$. The shared table is therefore not accidental: after burn-in, the defect sees the same sequence of background states.

Canonical orbit insufficiency (Fase 37). Generalizing the `F00` explanation to all representatives reveals a sharper constraint. Across the confirmed $T=15$ set there is only one canonical period-3 background orbit for `rule_73` and one for `rule_109`; this orbit identity is too coarse to determine the 13 shape families. The missing variable is not which orbit is reached, but how the localized defect is embedded into that orbit.

Embedding descriptor (Fase 38). Fase 38 tests descriptors based on the sampled orbit step, spatial rotation offset, IC start, IC length, defect anchor, and the first sampled defect state. The first sufficient descriptor is `(rule, sample_orbit_step, sample_rotation_offset, defect_state0)`. This closes the post-burn-in description: once the defect has crystallized into its first stable-cycle state, the family is determined. The descriptor is not yet a closed-form prediction from the initial condition, because `defect_state0` is measured after burn-in.

7.15 Pre-burn-in entry phase limits (Fase 39)

Fase 39 tests the remaining left side of the causal chain:

`(background, IC) -> burn-in -> defect_state0 -> family.`

For each of the 20 minimal $T=15$ representatives, the localized defect $D(t)=X(t) \text{ XOR } B(t)$ is traced every three ECA steps from $t=0$ to the sampling time $t=81$. The first time at which $D(t)$ belongs to one of the five stable cycle states is recorded as the entry time, together with the corresponding entry phase.

The entry is always fast: observed entry times are 0, 3, 6, 9, 12, with 15/20 representatives entering at $t=3$ and all 20 entering by $t=12$. However, fast entry is not the same as compact predictability. The descriptor tests include `rule+IC`, `rule+IC length`, post-hoc entry-time descriptors, and local background/IC windows of radius 1.5 around the initial perturbation. Exact predictors exist, but they are not compact: `rule+IC` uses 18 buckets with 17 singleton buckets, and the local windows produce 20 singleton buckets. Post-hoc descriptors involving entry time are exact for entry phase but do not predict it from the initial state.

The verdict is `NONCOMPACT_PREBURNIN_DESCRIPTOR_FOUND`. The stable-cycle entry point is reached after only a few applications of F^3 , but the tested pre-burn-in descriptors either fail or identify individual cases. Thus the current derivation is complete after burn-in and sharply delimited before burn-in: predicting `defect_state0` from the raw background/IC pair remains the irreducible part of the mechanism under the tested descriptor class.

7.16 Early causal-cone compression (Fase 40)

Fase 40 converts the negative result of Section 7.15 into a constructive causal compression test. Instead of seeking a closed-form pre-burn-in descriptor, it asks whether the relevant state can be recovered by simulating only the local causal cone of the initial perturbation.

For each of the 20 minimal $T=15$ representatives, the full reference remains the 256-cell simulation through $t=81$. The local predictor tests windows at $t = 3, 6, 9, 12$. Two window definitions are used: `span`, the full IC span plus a causal margin of t cells on each side, and `center`, the strict $2t+1$ window around the IC center. Boundary cells just outside the local window are supplied by the unperturbed periodic background orbit, which is valid because a radius-1 perturbation cannot reach beyond the cone within t steps.

The strict center cone succeeds at $t=12$. Its 25-cell, 12-step simulation matches the full-system defect state at $t=12$ in 20/20 representatives. When the detected stable-cycle phase is projected forward to $t=81$, it recovers `defect_state0` in 20/20 representatives. The compression ratio relative to the full 256-by-81 simulation is 69.1x. Earlier windows are incomplete: `center` at $t=9$ reaches 17/20 stable states, while `span` at $t=9$ reaches 19/20; both miss at least one representative.

The verdict is `EARLY_CONE_PREDICTOR_FOUND`. Fase 39 showed that no compact closed-form descriptor was found under the tested pre-burn-in descriptors. Fase 40 shows that the missing state is nevertheless locally causal: the mechanism does not require global information, only a short exact simulation of the radius-1 causal cone. The remaining formal problem is therefore not global dependence, but replacing this 25-cell, 12-step local computation with a symbolic account.

7.17 Minimal cone-table audit (Fase 41)

Fase 41 asks whether the successful 25-cell, 12-step cone from Section 7.16 contains a hidden smaller object: a sparse induced truth table, a reduced set of initial input bits, or a pruned causal subgraph. The analysis expands the cone into induced local updates of the form $\text{delta_f}(b,d)=f(b \text{ XOR } d) \text{ XOR } f(b)$, records the ordinary ECA truth-table entries used inside the cone, and computes the backwards dependency graph required to produce the final localized active defect support.

The result is deliberately negative at the table/input level. Across the 20 minimal representatives, the induced $(b,d) \rightarrow d_{\text{next}}$ tables are dense: they use 49.62 of the 64 possible local background/defect keys. All eight ordinary ECA truth-table entries are used in every representative. The final active defect support still depends on all 25 initial cone inputs. Thus there is no sparse truth-table shortcut and no input-bit elimination under this audit.

The only reduction is structural. Computing just the final active localized defect support uses 234..310 internal cone nodes, rather than all 325 nodes in the 25-cell-by-13-layer cone. This improves the circuit accounting but does not change the qualitative conclusion: the Fase 40 cone is close to minimal at the input level. The verdict is `STRUCTURAL_CONE_REDUCTION_ONLY`. The next symbolic target is therefore Boolean simplification of a dense 25-input, 12-step circuit, not further causal-support reduction.

7.18 ROBDD input-support audit (Fase 42)

Fase 42 applies reduced ordered binary decision diagrams (ROBDDs) to the dense 25-input, 12-step cone circuit. The goal is not to prove a globally minimal BDD over all possible variable orders, but to test the strongest simple Boolean reduction left open by Fase 41: whether any of the 25 cone inputs is semantically irrelevant to the active localized output functions.

For each of the 20 minimal `T=15` representatives, the Boolean circuit induced by the strict center cone is compiled into ROBDDs under the natural left-to-right cone variable order. The active localized outputs produce ROBDDs with 17,141..36,966 reachable nodes; the full 25-bit final vector produces 51,539..53,901 reachable nodes. In every representative, both the active-output functions and the full vector have support size 25/25.

The verdict is `BDD_NO_INPUT_REDUCTION`. ROBDD reduction therefore confirms the Fase 41 input-support result at the Boolean-function level: no initial cone variable can be eliminated without changing the represented output functions. BDD size remains order-dependent, so this is not a proof of global minimum BDD size over all `25!` orders. It does, however, rule out the key symbolic shortcut of Boolean input elimination. The remaining task is expression or BDD size reduction of a function that genuinely depends on all 25 inputs.

7.19 ROBDD order-sensitivity and targeted SIFT (Fase 43)

Fase 43 tests the natural follow-up to Section 7.18: if all 25 inputs are semantically necessary, perhaps the dense Boolean circuit is still compact under a better variable order. This section keeps the Fase 42 support result fixed and asks only about representation size.

Order-sensitivity preflight (Fase 43A). The three ROBDD orders already materialized in Fase 42 are compared: `natural`, `reverse`, and `center_out`. The best global order is `reverse`, but the improvement is small. Across the 20 minimal representatives, total active-output reachable nodes decrease from 552,476 under `natural` to 549,713 under `reverse`, a 0.5% reduction. The full 25-bit vector does not improve: 1,048,969 nodes under `natural` versus 1,049,085 under `reverse`. The `center_out` order is consistently poor, giving the worst active-output size in 20/20 representatives and a maximum order-sensitivity ratio of 3.119x. All tested orders preserve 25/25 active and vector support.

Targeted SIFT (Fase 43B). The most favorable representative from the preflight is `rule_73` on background `00111011`, family `F01`, with IC `00100100`. Its best Fase-42 active-output ROBDD is the `reverse` order with 16,061 reachable nodes. A one-pass variable-sifting search with checkpointing evaluates 580 orders on this representative. The best order found is `[22,21,20,19,23,24,18,17,16,15,14,13,12,11,10,9,8,7,6,5,4,3,2,1,0]`, with 16,056 active-output nodes and 52,698 full-vector nodes. This is a 5-node active-output reduction, or 0.031%, and remains far above the explicit 10,000-node compression gate. Support remains 25/25.

The verdict is not a global optimality proof over all `25!` orders, but it rules out the most direct representation shortcut tested here: simple variable reordering, even with a targeted SIFT pass on the best known candidate, does not make the dense cone compact. Any future symbolic shortcut must use a different representation or a different abstraction, not merely ROBDD variable ordering.

7.20 ANF degree audit of the T=15 cone (Fase 44)

Fase 44 tests a representation class independent from ROBDDs: algebraic normal form (ANF, or Zhegalkin polynomial) over GF(2). The question is whether the dense 25-input, 12-step causal cone has a compact polynomial representation even though BDD support and simple variable ordering do not simplify it.

The computation simulates exact truth tables in bit-packed form. Each active output has 2^{25} truth values; the simulation stores those values as `uint64` blocks, then unpacks one active output at a time and applies the Mobius transform to obtain ANF coefficients. This avoids stochastic sampling and does not rely on symbolic-regression heuristics.

Across the 20 minimal $T=15$ representatives, Fase 44 analyzes 174 active outputs. Active-output ANF degree ranges from 14 to 24; no output reaches the formal 25-variable ceiling. Active-output monomial counts range from 9,376 to 17,758,052. Sixty-seven of 174 active outputs have degree below 20, and every representative has at least one such lower-degree output.

The status is `LOW_OUTPUT_ANF_DEGREE_FOUND`. The cone is not algebraically uniform: some outputs remain extremely large as polynomials, but others have substantially lower degree. Fase 44 therefore turns the remaining symbolic problem from a uniform-density question into a stratification question: which outputs are algebraically simpler, and why?

7.21 ANF gradient laws (Fase 45)

Fase 45 analyzes the 174 active outputs from Fase 44 without new simulation. It finds that the ANF variation is not arbitrary; it is organized by distance from the cone center. Let `rel_pos` be the output position relative to the center cell of the 25-cell cone, and let $d = |\text{rel_pos}|$.

Degree gradient. For every active output,

`degree = 24 - d + epsilon`, with `epsilon` in `{0,1}`.

There are zero exceptions outside this one-bit epsilon band over all 174 active outputs. The center outputs ($d=0$) have degree exactly 24 in 13/13 cases. The immediate neighbors ($d=1$) have degree exactly 23 in 20/20 cases. The Pearson correlation between d and degree is -0.984525, and the linear fit is `degree ~ 24.093185 - 0.942523*d` with $R^2 = 0.969289$. The degree-24 cap also holds globally: no active output reaches degree 25.


Monomial-count gradient. Monomial counts obey an even cleaner spatial law:

`log10(monials) ~ 7.241925 - 0.307283*d`.

The Pearson correlation between d and `log10(monials)` is -0.999098, with $R^2 = 0.998197$. The slope differs from `-log10(2)` by only -0.006253, so the monomial count decays almost by a factor of two per cell away from the cone center.

ANF monomial gradient in the T=15 causal cone

$$\log_{10}(\text{monomials}) = 7.241925 - 0.307283 * d, R^2 = 0.998197$$



Blue points: mean by distance with min/max bars. Red line: fitted gradient.

Figure 4. ANF monomial gradient in the T=15 causal cone. Points show mean $\log_{10}(\text{monomials})$ by distance from the defect center with min/max bars; the red line is the fitted gradient.

The residual `epsilon` is not explained by a simple left/right symmetry. The center has `epsilon=0` in 13/13 cases; left outputs have `epsilon=1` in 32/87 cases, and right outputs in 26/74 cases. Rule identity also does not close the residual: `rule_73` has `epsilon=1` in 23/80 active outputs, while `rule_109` has `epsilon=1` in 35/94. Only 16/30 matched left/right output pairs have the same exact degree, so the defect is algebraically asymmetric even when its geometric presentation is visually paired.

The verdict is `ANF_GRADIENT_LAWS_CONFIRMED`. The T=15 defect acts as the focus of algebraic complexity in the cone. Complexity decreases linearly in ANF degree and almost exponentially in monomial count with distance from the defect center. This structure is orthogonal to the ROBDD results of Sections 7.18-7.19: BDD support and variable ordering do not simplify the function, but ANF exposes a spatial complexity gradient invisible to those tree-based audits.

7.22 Epsilon residual audit (Fase 46)

Fase 45 leaves one open bit in the ANF degree law. The backbone `degree = 24 - |rel_pos| + epsilon` has zero exceptions, but the residual `epsilon` is not explained by sign, rule identity, or simple left/right symmetry. Fase 46 asks whether this bit has a compact predictor from static features already available in the Fase 45 records.

The audit excludes `dist=0` and `dist=1`, where `epsilon=0` in all known cases. Including those rows would inflate accuracy without explaining the residual. The remaining dataset contains 141 active outputs across the 20 minimal $T=15$ representatives, with epsilon counts `{0: 83, 1: 58}` and a majority baseline of 58.87%.

Single-feature tests are evaluated by leave-one-representative-out validation. The best single feature is distance itself, at 64.89% mean accuracy. This is weak residual signal from the same variable that defines the main gradient, not a new explanatory law. The next strongest feature is `defect_phase` (64.53%), with phases 1 and 2 tending toward `epsilon=1`; this has plausible physical semantics but does not generalize strongly. `local_bg_3mer` reaches 64.18%. In contrast, `background_bit`, rule identity, and left/right sign collapse to the majority baseline, and `bg_transition` performs worse.

A depth-3 decision tree over all tested features reaches 73.05% training accuracy but only 55.65% mean leave-one-representative-out accuracy, with a 20.67% fold standard deviation. The tree's top feature is `family_id=F03`, followed by distance and local background features. This gap indicates representative/family memorization rather than a transferable epsilon rule.

The verdict is `EPSILON_REMAINS_RESIDUAL`. Fase 46 does not weaken the ANF gradient law; it separates the strong spatial backbone from a one-bit residual that is not captured by static rule, position, local-background, family, or cycle-phase descriptors. The remaining explanation, if one exists, likely requires dynamic features of the ANF computation rather than a static descriptor of the output position.

7.23 Dynamic ANF growth profile (Fase 47)

Fase 47 tests the explanation suggested by Fase 46: if `epsilon` is not predicted by static descriptors, it may be encoded in the temporal growth of the ANF itself. The computation reuses the exact bit-packed cone simulation of Fase 44, but instead of applying the Mobius transform only at the final layer, it computes ANF degree and monomial count for each active output at every layer $t = 1..12$.

The audit covers the same 20 minimal $T=15$ representatives and the same 141 nontrivial residual rows with `dist >= 2`. As an internal consistency check, the recomputed $t=12$ degrees and monomial counts are compared against Fase 44; there are zero mismatches.

The result is `EPSILON_DYNAMIC_RULE_FOUND`. The strongest single feature is `degree_growth_slope`, the linear slope of ANF degree over $t=1..12$. It reaches 98.58% training accuracy and 94.90% mean leave-one-representative-out accuracy with 9.44% fold standard deviation. This is a 30-point LORO improvement over the best static feature from Fase 46. `monomial_growth_slope` is weaker but still informative, with 73.37% LORO accuracy, and `t_first_full_degree` reaches 71.96%.

A depth-3 decision tree over the dynamic features reaches 86.52% training accuracy and 85.01% mean leave-one-representative-out accuracy. Its first split is `t_first_full_degree <= 11.5`, which predicts `epsilon=0`; outputs that do not reach their final degree before the last cone layer require additional monomial-growth and left/right-slope information. This gives a physical interpretation of the residual: `epsilon=1` is associated with outputs whose algebraic degree remains active near the causal horizon.

This is not a static pre-computation shortcut. The strongest feature uses the complete temporal degree trajectory through the final cone layer. The result is therefore best interpreted as a dynamic full-profile law of ANF growth: the epsilon bit is not visible in static output descriptors, but it is strongly encoded in how algebraic degree accumulates across the 12-step cone.

7.24 Early dynamic ANF horizon audit (Fase 48)

Fase 48 asks whether the Fase 47 dynamic predictor can be made earlier. The test reuses the stored $t=1..12$ ANF histories from Fase 47; no new ECA or cone simulation is performed. For each horizon K in `{6,8,9,10,11,12}`, it recomputes the future-blind feature `degree_growth_slope_K` using only the degree

profile from $t=1..K$, and evaluates it with the same 20-fold leave-one-representative-out protocol over the 141 nontrivial residual rows with $\text{dist} \geq 2$.

The result is `FULL_PROFILE_REQUIRED`. The horizon table for `degree_growth_slope_K` is:

Horizon K	LORO accuracy
6	61.74%
8	76.56%
9	75.27%
10	76.09%
11	79.47%
12	94.90%

The transition from $K=11$ to $K=12$ is not gradual: the final cone layer adds about 15 points of predictive accuracy. This supports the interpretation that `epsilon` is decided at the causal horizon. Outputs with `epsilon=1` are not merely high-complexity outputs; they are outputs whose ANF degree trajectory still carries decisive information in the last step of the 12-step cone.

The audit also tracks `t_first_full_degree_K`, but that feature uses the final expected degree from Fase 44 and is therefore not fully future-blind. The clean feature is `degree_growth_slope_K`. Under that feature class, no horizon before $K=12$ reaches the 90% gate. Thus Fase 47's law is not an early dynamic shortcut; it is a full-horizon profile law.

7.25 Generalization of the ANF gradient law (Fase 49)

Fase 49 tests whether the ANF gradient law discovered in Sections 7.20-7.21 is specific to the original length-8 $T=15$ representatives, or whether it also holds for the external $T=15$ backgrounds found in Fase 34. The external set contains eight replay-verified backgrounds: one length-9 `rule_73` background, five length-10 `rule_73` backgrounds, and two length-10 `rule_109` backgrounds. These are genuine external backgrounds, not rotations of the length-8 set.

The protocol reuses the exact bit-sliced Mobius ANF machinery from Fase 44: `WINDOW_CELLS=25` and `T_WINDOW=12`. Before ANF evaluation, each external background is replay-verified as $T=15$ using its minimal witness IC from Fase 34. All 8/8 backgrounds pass this replay gate. The Fase 34 witnesses have varying active defect widths, so the test does not assume constant visual defect width; the 25-cell cone is fixed by the radius-1, 12-step causal horizon.

Across the eight external backgrounds, Fase 49 analyzes 63 active outputs. The degree band remains exact:

$\text{degree} = 24 - d + \text{epsilon}$, with `epsilon` in $\{0,1\}$.

There are 0/63 exceptions. The active degree range is 16..24, and `epsilon` counts are $\{0: 39, 1: 24\}$.

The monomial-count law also generalizes quantitatively. The external fit is

$\log_{10}(\text{monomials}) \approx 7.224069 - 0.302890 \cdot d$, with $R^2 = 0.998263$.

Compared with the length-8 reference $\log_{10}(\text{monomials}) \approx 7.241925 - 0.307283 \cdot d$, the intercept differs by 0.25% and the slope magnitude by 1.43%. This is well inside the predefined generalization gate.

The verdict is `ANF_GRADIENT_GENERALIZES`. The ANF gradient is therefore not a length-8-specific artifact, at least across the original length-8 representatives and the external length-9/10 witnesses tested here. The question of whether this gradient is specific to the $T=15$ period or to a broader `rule_73/rule_109` mechanism family is addressed in Section 7.26.

7.26 Specificity of the ANF gradient beyond T=15 (Fases 50-53)

Fases 50--53 test whether the ANF gradient law from Fases 44--49 is a generic consequence of period, cone size, or active-support width, or whether it is tied to a more specific mechanism family.

Compact T=2 baselines (Fases 50--51). Fase 50 first tests the cleanest stationary local period-2 oscillator known from the quiescent-background catalog: `rule_108` on a zero background. Its active footprint is compact (`## <-> ###`), and at the 12-step comparison horizon the concrete active support has only 2 outputs in a single distance class. This is not enough support for a spatial active-output gradient. Fase 51 repeats the test for the four right-moving T=2 glider representatives `rule_20`, `rule_52`, `rule_148`, and `rule_180`, using the exact catalog ICs and a comoving final frame. The one-cell phase produces one active output and the two-cell phase control produces two adjacent active outputs, again in a single comoving distance class. The verdict for these compact T=2 baselines is `ANF_GRADIENT_T15_SPECIFIC`, but for a precise reason: compact active support prevents a meaningful gradient test. This does not by itself prove that the gradient is tied to period 15.

Wide periodic-background non-T15 cases (Fase 52). Fase 52 therefore moves to wide stationary oscillators over nontrivial periodic backgrounds. These cases have multiple active distance classes and are measured as XOR defects relative to the background orbit, matching the T=15 convention. The key positive witness is `rule_109` on background `1011`, with `T_local=10`. At the common 12-step horizon it gives

`log10(monials) ~ 7.258688 - 0.307674*d`, with `R^2 = 0.999349`.

The slope differs from the T=15 reference `-0.307283` by only 0.13%. This is not merely a similar trend; within this finite audit it is the same spatial monomial-decay gradient. Other non-T15 periodic-background cases are weaker or ambiguous: `rule_73` on background `0010` with `T_local=10` has too few active distance classes at the 12-step horizon (`reliable=no`); `rule_73` on background `0011` with `T_local=12` gives slope `-0.296180` but only `R^2 = 0.877281`; and `rule_94` on background `0010` with `T_local=3` is flat at its own period (slope `0.000523`, `R^2 = 0.049180`) but develops a slope `-0.341994` with `R^2 = 0.955626` only when oversampled to `T_WINDOW=12`. That last case is treated as a horizon effect rather than the same natural period law. The Fase 52 verdict is `ANF_GRADIENT_MECHANISM_DEPENDENT`.

External-family test (Fase 53). Fase 53 then tests wide candidates outside the `rule_73/rule_109` family, using the shortest catalog witnesses at maximum support `span=11`. `rule_54` on background `0010` with `T_local=4` has slope `-0.002354` at its own period: effectively flat. `rule_94` on background `0001` with `T_local=6` has slope `0.000000` and `R^2 = 0.000000` at its own period. `rule_133` on background `1011` with `T_local=6` also has slope `0.000000` and `R^2 = 0.000000` at its own period. These own-period tests have enough active support to be informative; the flat result is not a support artifact. At the common 12-step horizon, some external cases acquire nonzero slopes, but they do not reproduce the T=15-quality combination of support reliability, slope, and `R^2`. The Fase 53 verdict is `ANF_GRADIENT_FAMILY_73_109`.

Interpretation. The ANF gradient is not a consequence of period, cone size, or active-support width alone. It is mechanism-dependent. The confirmed non-T15 witness in this audit is the `rule_109/background 1011/T=10` case, while compact T=2 oscillators and the external families tested (`rule_54`, `rule_94`, and `rule_133`) do not reproduce the T15-quality gradient at their natural periods. Section 7.27 tests whether this non-T15 witness broadens across additional `rule_73/rule_109` family members. No claim is made about untested ECA families.

7.27 Additional family robustness test for the ANF gradient (Fase 54)

Fase 54 tests whether the non-T15 gradient witness from Section 7.26 is an isolated case or part of a broader natural-period law within the `rule_73/rule_109` periodic-background family. The test uses the same 25-input bit-sliced Mobius ANF engine as Fases 52--53. The primary criterion is the natural-period horizon `T_WINDOW=T_local`; the common 12-step horizon is reported only as a secondary comparison.

Three additional family witnesses are selected from the periodic-background catalog:

- rule_109, background 1011, T_local=6, IC 00001001;
- rule_109, background 1101, T_local=10, IC 0001000;
- rule_73, background 0010, T_local=6, IC 1100111.

At their natural periods, none reproduces the T15-quality gradient. The two T=6 cases are essentially flat: rule_109/1011 has slope 0.000026 and $R^2 = 0.604938$, while rule_73/0010 has slope -0.000027 and $R^2 = 0.604938$. The rule_109/1101/T=10 case has a nonzero slope (-0.209698) but lower fit quality ($R^2 = 0.880488$) and a 31.76% deviation from the T15 reference slope, so it is not comparable to the T15 law.

At the common 12-step horizon, two of the new cases do show T15-like slopes: rule_109/1011/T=6 gives slope -0.303174 with $R^2 = 0.999487$, and rule_73/0010/T=6 gives slope -0.320463 with $R^2 = 0.999687$. Because these gradients appear only after oversampling to the 12-step horizon, they are classified as horizon effects rather than natural-period witnesses. The Fase 54 verdict is therefore

ANF_GRADIENT_ISOLATED_WITNESS under this targeted three-case test: the strong non-T15 natural-period witness remains the rule_109/background 1011/T=10 case from Fase 52. Section 7.28 then replaces this local conclusion with a full catalog census.

7.28 ANF gradient census across the periodic-background catalog (Fase 55)

Fase 55 turns the targeted robustness question into a catalog-level census. The test scans all stationary periodic-background oscillator groups with `span >= 11`, excluding compact T_local=2 baselines and the original T_local=15 family. One IC is evaluated per (rule, background, T_local) group: groups already tested in Fases 52--54 keep their exact previous ICs as consistency baselines, while new groups use maximum span and then shortest word as the tie-breaker.

The preflight census contains 66 groups across six rules:

- period distribution {3: 18, 4: 8, 6: 22, 8: 6, 10: 8, 12: 4};
- rule distribution {54: 4, 73: 17, 94: 12, 109: 17, 133: 12, 147: 4};
- seven groups marked already_tested=true.

The ANF engine performs 128 measurements, with zero packed/concrete discrepancies. The category counts are:

Category	Count
NATURAL_PERIOD_STRONG	2
HORIZON_ACCEPTABLE	3
HORIZON_ARTIFACT	20
INSUFFICIENT_SUPPORT	3
NEGATIVE	38

Among previously untested cases, the census finds two NATURAL_PERIOD_STRONG witnesses, both in rule_109:

Case	Slope	R^2	Delta vs T15
rule_109/background 0011/T=12/IC 10010100	-0.298274	0.998341	2.93%
rule_109/background 1100/T=12/IC 00101001	-0.298274	0.998341	2.93%

Because these cases have T_local=12, their natural-period measurement and the common T_WINDOW=12 measurement are the same run. They are valid natural-period witnesses, but they are not two independent confirmations across different horizons.

The census also finds three `HORIZON_ACCEPTABLE` cases: the already-tested baseline `rule_109/background 1011/T=10` and two new `rule_109/T=8` witnesses:

Case	Natural-period fit	Common-horizon fit
<code>rule_109/background 0110/T=8/IC 0000011</code>	slope <code>-0.106802</code> , R^2 <code>0.617294</code>	slope <code>-0.298928</code> , R^2 <code>0.998276</code> , delta <code>2.72%</code>
<code>rule_109/background 1100/T=8/IC 00000110</code>	slope <code>-0.106802</code> , R^2 <code>0.617294</code>	slope <code>-0.298928</code> , R^2 <code>0.998276</code> , delta <code>2.72%</code>
<code>rule_109/background 1011/T=10/IC 00000001</code>	slope <code>-0.196127</code> , R^2 <code>0.922575</code>	slope <code>-0.307674</code> , R^2 <code>0.999349</code> , delta <code>0.13%</code>

The negative pattern is equally important. No `rule_73` case reaches `NATURAL_PERIOD_STRONG` or `HORIZON_ACCEPTABLE`. The external rules tested (`rule_54`, `rule_94`, `rule_133`, and `rule_147`) also produce no strong or acceptable witness. Many `T_local <= 6` cases are `HORIZON_ARTIFACT`: they develop T15-like slopes only after oversampling to the 12-step horizon, not at their natural period.

The Fase 55 verdict is `NEW_NATURAL_PERIOD_WITNESS_FOUND`. This updates the interpretation from Sections 7.26--7.27: the ANF gradient is not merely an isolated non-T15 witness, but neither is it shared symmetrically by `rule_73/rule_109` as a family. In the censused catalog, the robust non-T15 evidence is concentrated in `rule_109`.

7.29 Structural audit of the `rule_109` concentration (Fase 56)

Fase 56 asks why the Fase 55 witnesses concentrate in `rule_109`. It does not run new ANF simulations; instead it audits the Fase 55 census by comparing background rotation/complement structure, the rule-level algebraic normal forms of `rule_73` and `rule_109`, and direct cross-rule cases on shared backgrounds.

The resulting statuses are:

- `orbit_symmetry_status = PARTIAL`;
- `rule_anf_status = RULE109_CENTER_MEDIATED_CONFIRMED`;
- `cross_rule_status = RULE109_SPECIFIC_ON_SHARED_BACKGROUNDS`;
- `overall_status = RULE109_SYMMETRY_MECHANISM_CANDIDATE`.

The cyclic-orbit evidence is informative but incomplete. Several positive `rule_109` witnesses belong to the rotation orbit of background `0011`, but the already-confirmed `rule_109/background 1011/T=10` witness belongs to a different rotation orbit. Complement evidence is also partial: for example, backgrounds `0011` and `1100` are complementary and both are `NATURAL_PERIOD_STRONG` at `T=12`, while other complements are absent from the census or negative at the tested support threshold.

The rule-level ANF contrast is exact:

Rule	ANF over (L,C,R)	Center alone?	LR without center?
73	<code>1 XOR L XOR C XOR R XOR LR XOR LCR</code>	yes	yes
109	<code>1 XOR L XOR LC XOR R XOR CR XOR LCR</code>	no	no

Thus, in `rule_109` the center cell contributes only through neighbor-mediated terms (`LC`, `CR`, `LCR`); there is no isolated center term and no `LR` term that bypasses the center. In `rule_73`, the center contributes directly through `C`, and the neighbors also interact through `LR` without the center. This is a deterministic algebraic distinction, not a statistical fit.

The direct cross-rule comparisons are consistent with the rule-specific interpretation:

<code>rule_109</code> positive case	Matching <code>rule_73</code> case
<code>bg=0011/T=12/IC 10010100 -> NATURAL_PERIOD_STRONG</code>	<code>bg=0011/T=12/IC 10001010 -> NEGATIVE</code>

rule_109 positive case	Matching rule_73 case
bg=0110/T=8/IC 0000011 -> HORIZON_ACCEPTABLE	absent from census
bg=1011/T=10/IC 00000001 -> HORIZON_ACCEPTABLE	absent from census
bg=1100/T=8/IC 00000110 -> HORIZON_ACCEPTABLE	bg=1100/T=8/IC 0011111 -> NEGATIVE
bg=1100/T=12/IC 00101001 -> NATURAL_PERIOD_STRONG	bg=1100/T=12/IC 00000011 -> NEGATIVE

Fase 56 therefore identifies a mechanism candidate, not a closed proof. Background orbit structure matters, but it does not fully explain the witness set. The sharper invariant is the rule-level algebraic contrast: rule_109 mediates center dependence through neighbor interactions, and the available cross-rule comparisons show that matching rule_73 cases do not become positive witnesses.

7.30 Center-mediated causality audit (Fase 57)

Fase 57 tests whether the center-mediated structure identified in Fase 56 is necessary and/or sufficient for the ANF-gradient witnesses found in the Fase 55 catalog. It does not run new ECA or cone simulations. Instead, it computes the 3-variable ANF of all 256 ECA rules and joins the resulting rule-level classification to the existing catalog census.

The local-rule classification uses:

- C_alone: the isolated C monomial appears;
- LR_no_center: the LR monomial appears without the center;
- center_mediated: C_alone=False and LR_no_center=False;
- strict_center_mediated: center_mediated=True and at least one center interaction monomial exists.

Across all 256 ECA rules, 64 rules satisfy center_mediated=True, and 56 of those are strict center-mediated rules. Within the Fase 55 catalog, the tested rules are 54, 73, 94, 109, 133, and 147. Only rule_109 is center-mediated under the above definition.

The catalog-level table is:

Rule	ANF over (L,C,R)	Center-mediated?	Cases	Positive witnesses
54	L XOR C XOR R XOR LR	no	4	0
73	1 XOR L XOR C XOR R XOR LR XOR LCR	no	17	0
94	L XOR C XOR LC XOR R XOR CR XOR LCR	no	12	0
109	1 XOR L XOR LC XOR R XOR CR XOR LCR	yes	17	5
133	1 XOR L XOR R XOR LR XOR LCR	no	12	0
147	1 XOR C XOR LR	no	4	0

The necessity test is clean: the catalog contains 5 positive witnesses (NATURAL_PERIOD_STRONG or HORIZON_ACCEPTABLE), and 0/5 occur in rules with center_mediated=False. Thus center mediation is necessary for the observed positive witnesses inside the Fase 55 catalog.

The sufficiency test fails: the catalog contains 12 non-positive center-mediated cases, all within rule_109. These include negative low-period cases such as rule_109/bg=0011/T=3, rule_109/bg=0011/T=6, rule_109/bg=0110/T=3, and rule_109/bg=1100/T=6, as well as horizon-artifact cases such as rule_109/bg=1011/T=6 and rule_109/bg=1101/T=6.

The resulting status is CAUSAL_CANDIDATE_NECESSARY_NOT_SUFFICIENT. The Fase 56 candidate therefore survives as a necessary structural condition within the catalog, but not as a complete causal explanation. The

next discriminant is not rule-level center mediation alone, but center mediation plus period/horizon structure. A closed causal proof would require an intervention step, such as constructing synthetic rules or rule variants that preserve or break the center-mediated terms while holding the catalog protocol fixed.

7.31 Period/horizon audit within rule_109 (Fase 58)

Fase 58 asks whether the missing second condition from Fase 57 is simply the period/horizon relation. The analysis restricts the Fase 55 census to the 17 rule_109 cases, all of which share the center-mediated local ANF structure, and tests whether T_{local} , the common horizon $T_{WINDOW}=12$, or the oversampling ratio $12/T_{local}$ separates the 5 positive witnesses from the 12 non-positive cases. No new ECA or ANF simulation is run.

The rule_109 case distribution is:

Background	T_{local}	$12/T_{local}$	Category
0011	3	4.000	NEGATIVE
0011	6	2.000	NEGATIVE
0011	8	1.500	NEGATIVE
0011	10	1.200	NEGATIVE
0011	12	1.000	NATURAL_PERIOD_STRONG
0110	3	4.000	NEGATIVE
0110	6	2.000	NEGATIVE
0110	8	1.500	HORIZON_ACCEPTABLE
1011	6	2.000	HORIZON_ARTIFACT
1011	10	1.200	HORIZON_ACCEPTABLE
1100	3	4.000	NEGATIVE
1100	6	2.000	NEGATIVE
1100	8	1.500	HORIZON_ACCEPTABLE
1100	10	1.200	NEGATIVE
1100	12	1.000	NATURAL_PERIOD_STRONG
1101	6	2.000	HORIZON_ARTIFACT
1101	10	1.200	NEGATIVE

Period and horizon are informative. The cleanest high-precision rule, $T_{local} == 12$ (equivalently $12/T_{local} == 1$), has no false positives but captures only 2/5 positives. Conversely, $T_{local} \geq 8$ (equivalently $12/T_{local} \leq 1.5$) captures all 5 positives but also admits 4 false positives: $bg=0011/T=8$, $bg=0011/T=10$, $bg=1100/T=10$, and $bg=1101/T=10$.

Thus, no period/horizon-only rule separates the 17 rule_109 cases perfectly. The resulting status is PERIOD_HORIZON_PARTIAL_DISCRIMINANT.

The key residual is rule_109/bg=0011/T=8: it satisfies the horizon threshold $T_{local} \geq 8$ and $12/T_{local} \leq 1.5$, yet remains NEGATIVE, whereas $bg=0110/T=8$ and $bg=1100/T=8$ are HORIZON_ACCEPTABLE. These three backgrounds lie in the same cyclic rotation orbit, so the residual is not explained by the coarse background orbit class. Fase 58 therefore narrows the next causal layer to background phase, IC placement, or alignment inside the oscillator mechanism.

7.32 IC/background alignment audit and validation (Fases 59--60)

Fase 59 audits the residual left by Fase 58. The three rule_109/T=8 cases share center mediation and the same period/horizon threshold ($T_{local}=8$, $12/T_{local}=1.5$), but split into one negative and two acceptable-horizon witnesses:

Background	Category	IC word	IC active bits	IC offsets mod 4	ic_span
0011	NEGATIVE	1000010	(0, 5)	(0, 1)	6
0110	HORIZON_ACCEPTABLE	0000011	(5, 6)	(1, 2)	2
1100	HORIZON_ACCEPTABLE	00000110	(5, 6)	(1, 2)	2

This gives a clean local discriminator. The two positive $T=8$ cases use adjacent active IC bits at offsets (1,2) with span 2, while the negative case uses separated active bits at offsets (0,1) with span 6. The Fase 59 status is therefore `ALIGNMENT_DISCRIMINANT_FOUND` for the $T=8$ residual.

Background-subtracted descriptors also carry signal. For the same three cases, `xor_defect`, `defect_phase_offset`, and defect weights separate the negative value from the positive values. However, the two positives do not share one identical defect value. These descriptors therefore provide partial alignment evidence, not an exact rule.

Fase 60 then validates the Fase 59 discriminator against all 17 `rule_109` cases in the Fase 55 census. The local rule

```
ic_span == 2 and ic_active_offsets_mod4 == (1,2)
```

is perfect on $T=8$ alone (TP=2, FP=0, TN=1, FN=0, accuracy 1.000), but it captures only 2/5 positives across all `rule_109` cases (TP=2, FP=0, TN=12, FN=3, accuracy 0.824). In particular, `ic_span` is not a global discriminator: positive and non-positive cases both occur with `ic_span` values 1, 2, and 6.

The only descriptor that perfectly separates the 17 catalog cases is the exact `ic_active_bits` tuple. This is useful as an audit result, but it is lookup-like: it identifies the selected IC word pattern rather than a compact causal rule. Fase 60 therefore assigns the status `ALIGNMENT_LOOKUP_ONLY`.

The causal audit has reached the limit of static rule, period, background orbit, and IC-placement descriptors available in the catalog. The chain is now: center mediation is necessary but not sufficient; period/horizon is informative but incomplete; IC alignment explains the $T=8$ residual locally but does not become a compact global condition. A stronger causal explanation would require dynamic alignment features or an intervention experiment rather than further static descriptors alone.

7.33 Dynamic alignment audit and minimal union (Fases 61--62)

Fase 61 tests whether the remaining causal signal is dynamic rather than static. For each of the 17 `rule_109` cases in the Fase 55 census, the analysis simulates two systems over the common horizon $t=1..12$: the IC word superposed on its periodic background, and the pure periodic background alone. The measured defect is the XOR difference

```
defect(t) = state_with_IC(t) XOR background_only(t).
```

This avoids reconstructing the background from `T_local`: the background is itself evolved under `rule_109`, in parallel with the IC state. The simulation uses a wide window (`WIDTH=256`) and places the IC with phase compatible with the background period. For each time step, the audit records defect size, span, compactness, center of mass, and derived growth/drift descriptors.

No single dynamic descriptor separates all five positive witnesses from the twelve non-positive cases. However, several thresholds give no false positives:

Descriptor rule	TP	FP	TN	FN	Captured positives
<code>max_defect_size >= 12</code>	1	0	12	4	<code>bg=1011/T=10/word=0000001</code>

Descriptor rule	TP	FP	TN	FN	Captured positives
center_drift_abs <= 0.0	2	0	12	3	bg=0110/T=8/word=0000011; bg=1011/T=10/word=0000001
size_growth_total <= -3	2	0	12	3	bg=0011/T=12/word=10010100; bg=1100/T=12/word=00101001

The Fase 61 status is therefore **DYNAMIC_PARTIAL**: dynamic descriptors carry high-precision signal, but as subfamily-specific signatures rather than a single separating law.

Fase 62 tests one predeclared minimal union, not a broad combinatorial search:

size_growth_total <= -3 OR center_drift_abs <= 0.0.

The result is:

Rule	TP	FP	TN	FN	Accuracy	Precision	Recall
size_growth_total <= -3 OR center_drift_abs <= 0.0	4	0	12	1	0.941	1.000	0.800

The union captures four positive witnesses with no false positives: bg=0011/T=12/word=10010100, bg=0110/T=8/word=0000011, bg=1011/T=10/word=0000001, and bg=1100/T=12/word=00101001. The single remaining false negative is bg=1100/T=8/word=00000110. Adding max_defect_size >= 12 does not improve recall, because it captures a positive already covered by center_drift_abs <= 0.0.

The Fase 62 status is **DYNAMIC_UNION_PARTIAL**. This result is stronger than noise: the union has precision 1.000 and no false positives. But it still does not close the causal explanation. The case bg=1100/T=8/word=00000110 survives center-mediated structure, period/horizon thresholds, static IC-alignment descriptors, single dynamic descriptors, and the minimal dynamic union. It is therefore the most informative remaining target for a future intervention experiment.

7.34 Algebraic intervention preflight (Fase 63)

Fase 63 tests whether the residual identified by Fase 62 can be probed by minimal algebraic interventions on the local ANF of rule_109. The local rule is

rule_109 = 1 XOR L XOR LC XOR R XOR CR XOR LCR.

Four single-monomial edits are tested:

Intervention	Synthetic rule	Local ANF	Center mediated?	Fase 55 census?
remove LC	173	1 XOR L XOR R XOR CR XOR LCR	yes	no
remove CR	229	1 XOR L XOR LC XOR R XOR LCR	yes	no
add isolated C	161	1 XOR L XOR C XOR LC XOR R XOR CR XOR LCR	no	no
add LR without center	205	1 XOR L XOR LC XOR LR XOR R XOR CR XOR LCR	no	no

The ANF edits are verified exactly. The next question is whether any synthetic rule preserves the stationary oscillator support needed to ask the ANF-gradient question on the residual background bg=1100. Fase 63 therefore runs the same periodic-background oscillator preflight as earlier phases, restricted to bg=1100 and IC words of length 1..8.

The result is:

Synthetic rule	Processed ICs	Stationary hits	Moving hits	Period-1 aliases
173	502	0	0	12
229	502	0	0	8
161	502	0	0	0
205	502	0	0	0

No synthetic rule has a stationary oscillator on `bg=1100` in this preflight. Consequently, no cone-ANF measurement is run: there is no comparable stationary witness on the residual background. The Fase 63 status is `ALGEBRAIC_INTERVENTION_PREFLIGHT_BLOCKED`.

This is an informative block rather than a failed measurement. At the monomial-edit resolution tested here, the `rule_109` mechanism is not separable into independently removable local ANF terms while preserving the oscillatory support. The center-mediated terms are entangled with viability of the oscillator itself. A causal intervention must therefore be finer or more conditioned than a whole-monomial ANF edit, for example by moving to Hamming-1 truth-table neighbors or by targeting the surviving residual with a more local dynamic analysis.

7.35 Hamming-1 neighbourhood audit (Fase 64)

Fase 64 tests whether the Fase 63 block is merely an artifact of intervention granularity. Whole-monomial ANF edits can flip more than one truth-table bit at once. The finest possible ECA intervention is therefore a Hamming-1 neighbour:

`rule_i = 109 XOR (1 << i)`, for `i = 0..7`.

The resulting neighbours are `108`, `111`, `105`, `101`, `125`, `77`, `45`, and `237`. Note that the bit-4 neighbour is `125`, not `93`, since `109 XOR 16 = 125`.

Each neighbour is swept on `bg=1100` with IC words of length `1..8` under the same preflight protocol as Fase 63:

Bit flipped	Rule	Local ANF	Stationary hits	Moving hits	Max span	Periods
0	108	C XOR LR	237	0	9	[2]
1	111	1 XOR L XOR LR XOR LC	0	0	0	[]
2	105	1 XOR R XOR C XOR L	0	0	0	[]
3	101	1 XOR R XOR L XOR LC	0	0	0	[]
4	125	1 XOR R XOR CR XOR LR	0	0	0	[]
5	77	1 XOR R XOR CR XOR L XOR LR XOR LC	0	0	0	[]
6	45	1 XOR R XOR CR XOR L	0	0	0	[]
7	237	1 XOR R XOR CR XOR L XOR LC	0	0	0	[]

Seven of eight neighbours therefore behave like the Fase 63 monomial edits: they do not preserve stationary or moving oscillator support on `bg=1100`. The single survivor is `rule_108`, produced by flipping bit 0. It yields 237 stationary witnesses, but all are compact period-2 oscillators with maximum span 9. The strongest measured witness is `rule_108/bg=1100/T=2/word=00000001`.

Its ANF measurement is not comparable to the `rule_109/T=8` residual:

T_WINDOW	Active outputs	Distance classes	Slope	R^2	Comparable
2	5	4	-0.017102	0.264706	false
12	5	4	-0.017102	0.264706	false

Thus, Hamming-1 perturbation does not produce a comparable modified `rule_109/bg=1100/T=8` mechanism. It either destroys the oscillator support entirely, or in the bit-0 case collapses into the already known compact `rule_108/T=2` attractor. The Fase 64 status is `HAMMING1_WITNESSES_FOUND_NOT_COMPARABLE`.

Together, Fases 63--64 establish a local perturbation boundary around the residual mechanism. Neither monomial-level ANF edits nor atomic truth-table bit flips yield a comparable stationary oscillator with an ANF-gradient signature on `bg=1100`. The `rule_109/bg=1100/T=8` residual is therefore locally isolated under the tested ECA perturbations; a stronger causal experiment would need a conditioned intervention that preserves the orbit rather than changing the global rule table uniformly.

7.36 Context-frequency and persistence audit of the residual (Fases 65--66)

Fase 65 asks whether the remaining residual can be distinguished by the local contexts actually used by its active defect cells. The test compares the residual `rule_109/bg=1100/T=8/word=00000110` with the other four positive `rule_109` witnesses and three `bg=1100` negative controls. Each case is evolved to `t=50` with width 256. As in the dynamic alignment audit, the pure background is evolved in parallel and the defect is measured as

`defect(t) = state_with_IC(t) XOR background_only(t)`.

For each active defect cell, the local context index is

`(L << 2) | (C << 1) | R`,

so the eight possible contexts are `000` through `111`. For `t>0`, contexts are read from `state_with_IC(t-1)`; for `t=0`, they are read from the pure background frame.

The aggregate context-frequency test is negative:

- contexts present in every positive and no negative control: none;
- contexts used by the residual and by no other positive: none;
- contexts used by every other positive but absent from the residual: none;
- contexts never used by any selected defect cell: none.

Thus, the Fase 65 status is `CONTEXT_UNDISCRIMINATED`, with `NO_UNUSED_CONTEXT_INTERVENTION_CANDIDATE`. The selected cases exercise all eight contexts, so no uniform context flip is obviously safe.

The most informative number is the L1 context-frequency neighbourhood of the residual:

Rank	Case	Category	L1 distance from residual
1	<code>bg=1100/T=10/word=00111001</code>	<code>NEGATIVE</code>	0.104
2	<code>bg=1100/T=12/word=00101001</code>	<code>NATURAL_PERIOD_STRONG</code>	0.256
3	<code>bg=1011/T=10/word=00000001</code>	<code>HORIZON_ACCEPTABLE</code>	0.503
4	<code>bg=0011/T=12/word=10010100</code>	<code>NATURAL_PERIOD_STRONG</code>	0.618
5	<code>bg=1100/T=6/word=00100110</code>	<code>NEGATIVE</code>	0.671
6	<code>bg=1100/T=3/word=00001110</code>	<code>NEGATIVE</code>	0.802
7	<code>bg=0110/T=8/word=0000011</code>	<code>HORIZON_ACCEPTABLE</code>	1.174

This creates a sharper paradox: the positive residual is closer in aggregate context usage to a negative $bg=1100/T=10$ case than to any other positive witness.

Fase 66 tests the cheapest explanation for that paradox: perhaps the residual is not a genuine oscillator, but a long transient that only passed the finite `HORIZON_ACCEPTABLE` threshold used in the earlier census. The same residual is evolved to $t=500$, again measuring the background-subtracted defect. The result refutes the transient hypothesis:

- classification: `PERSISTENT_OSCILLATOR`;
- status: `RESIDUAL_CONFIRMED_PERSISTENT`;
- collapse step: none;
- exact period observed in the last 100 steps: 8;
- center slope in the last 100 steps: $5.591147350029335e-05$;
- final defect size: 6;
- tail defect-size range: 5..8.

Therefore, the residual is not a false positive caused by a pre-collapse window. It is a genuine period-8 oscillator whose aggregate context frequencies do not explain why it remains positive while the nearby $bg=1100/T=10$ control is negative. The discrimination is not in the set or aggregate frequency of contexts; it must lie in their phase ordering, trajectory, or another time-resolved invariant.

7.37 Phase organization and compressed causal-complexity audit (Fases 67--68)

Fase 67 asks whether the residual differs from its nearest negative control when the defect trajectory is decomposed by internal phase rather than aggregated over time. The residual `rule_109/bg=1100/T=8/word=00000110` is compared with the nearest context-frequency negative from Fase 65, `rule_109/bg=1100/T=10/word=00111001`.

Both defects are genuine periodic trajectories. The residual is verified with exact period 8 after burn-in, and the negative control is verified with exact period 10. Thus the comparison is not oscillator versus transient; it is a comparison between two stable periodic circuits.

The phase-level comparison reveals partial structure:

- residual phases 1 and 4 have context L1 distance 0.000 to their best matching negative phases, so the two oscillators share part of their local circuit structure;
- residual phase 0 has best context L1 distance 0.571 and dominant contexts `011` and `111`;
- residual phase 3 has best context L1 distance 0.524, active Jaccard 0.182, and a uniquely dominant `100` context;
- residual phase 7 has best context L1 distance 0.400 and dominant context `111`.

The Fase 67 status is therefore `PHASE_DISCRIMINANT_FOUND` for the pairwise residual-versus-nearest-negative comparison. The discrimination is not the set of contexts used by the defect, but their organization by phase.

Fase 67b cross-validates the phase-dominant contexts against the five positive `rule_109` witnesses and three $bg=1100$ negative controls. The primary context is `100`, because it was uniquely dominant in residual phase 3, the phase with strongest spatial separation from the negative cycle. The result is consistent but not separating:

Context	Positives dominant	Positives unique	Negatives dominant	Negatives unique
<code>011</code>	3/5	1/5	2/3	0/3
<code>100</code>	4/5	1/5	2/3	0/3
<code>111</code>	4/5	4/5	3/3	2/3

Thus `100` replicates across positives, but it is not specific to positives. The Fase 67b status is `PHASE_CROSSVAL_CONSISTENT`, not a clean causal separator.

Fase 68 then tests a compressed causal-state proxy over all 17 `rule_109` catalog cases. Each time step from `t=0..100` is symbolized as

```
(dominant_context(t), defect_size_bucket(t)),
```

where the size bucket is one of `0`, `1-3`, `4-6`, `7-9`, or `10+`. The resulting symbolic sequence is measured using bigram entropy, unique transitions, Lempel-Ziv complexity, detected symbolic period, and number of unique symbols.

No compressed complexity metric separates the 5 positive witnesses from the 12 non-positive cases:

- perfect rules: 0;
- perfect non-period complexity rules: 0;
- best no-false-positive rule: `period_detected >= 12`, with TP=2, FP=0, TN=12, FN=3, precision 1.000, recall 0.400;
- best non-period rule: `bigram_entropy >= 3.673269689515108`, with TP=2, FP=1, TN=11, FN=3, precision 0.667, recall 0.400.

The best no-false-positive rule is only the symbolic period threshold, which recapitulates the earlier period/horizon result rather than adding a new causal-state signal. The non-period complexity metrics overlap strongly between positives and negatives. For example, positive `bigram_entropy` ranges from 0.322 to 3.673, while negative `bigram_entropy` ranges from 0.081 to 4.040.

The combined status is `PHASE_PARTIAL_NOT_SEPARATING`. The residual is not explained by any scalar temporal descriptor tested: dynamic summaries, global perturbations, aggregate contexts, phase-dominant contexts, and compressed causal-complexity proxies all fail to provide a compact separator. The remaining information appears to reside in richer spatial or spatiotemporal cone structure, not in compressed temporal summaries of the defect. This closes the scalar causal audit around the `rule_109` residual and motivates future work based on full spatial snapshots, richer causal-state reconstruction, or external review before adding further internal phases.

7.38 External `rule_73` witnesses and partial horizon robustness (Fases 77--79)

Fases 77--79 reopen external validation only after fixing the Fase 55 comparison predicate and candidate-selection rules. No threshold is fitted to the new cases. Fase 77 first evaluates 76 natural-period groups from four rules not present in the Fase 55 catalog (`62`, `118`, `131`, and `145`) over primitive length-8 backgrounds. All 76 candidates have natural period 3, all fits are reliable, and none is comparable to the T15 ANF gradient. The best observed fit has $R^2=0.619571$, well below the predeclared `0.95` threshold. This is a bounded negative result for the available period-3 holdout, not a test of external rules at periods 8--12.

Fase 78 then queries the completed length-8 physical sweep for previously unmeasured `rule_73` oscillators in that longer-period range. The fixed holdout contains 25 stationary groups on 18 primitive backgrounds:

Natural period	Cases	Comparable witnesses
8	1	0
10	6	0
12	18	9

The nine witnesses occur on nine distinct backgrounds and are measured only at their natural period. The strongest fit has $R^2=0.999252$; the closest slope is `-0.307370`, a 0.03% deviation from the T15 reference. An independent no-checkpoint remeasurement reproduces the packed result, and no packed versus concrete

discrepancy is observed. These are the first natural-period non-`rule_109` witnesses in the two completed catalogs. Because `rule_73` has an isolated `C` term and an `LR` term without the center, it is not center-mediated under the Fase 57 definition. Center mediation is therefore not necessary beyond the original Fase 55 catalog; the earlier necessity result remains valid only within its stated catalog.

All nine new witnesses have `T_local=12`, which coincides with the common horizon resonance identified in Fase 76. Fase 79 therefore remeasures all 18 `rule_73/T=12` holdout cases at predeclared neighboring horizons 10, 14, and 16. The nine natural-period witnesses are the primary cohort, and the nine natural-period negatives are controls:

Cohort	Cases	Comparable at h10	Comparable at h14	Comparable at h16	Survive any neighbor
Fase 78 witnesses	9	1	2	0	3
Fase 78 controls	9	0	0	0	0

Each of the three survivors is comparable at one neighboring horizon only; no case survives at horizon 16. Across the 27 control measurements, none crosses the unchanged comparison threshold. Thus the Fase 78 result is not a pure point resonance at horizon 12, but neither is it horizon-invariant.

The combined status is `RULE73_LEN8_NEIGHBOR_HORIZON_PARTIAL`. The ANF gradient has a first natural witness outside `rule_109` and finite neighboring-horizon robustness in a subset of backgrounds. The claim is strictly background-conditioned: it covers one additional rule, primitive length-8 backgrounds, nine natural-period witnesses, and three survivors at neighboring horizons. It does not establish universal generalization across ECA rules or backgrounds.

7.39 Global audit of the historical period cap (Fases 89--90)

The periodic-background catalogs used a source detector restricted to exact periods `2..16`. Fase 88 exposed one concrete consequence: 40 trajectories previously labelled nonstationary were stationary with period 30 when checked to `t=1000`. Fase 89 first audited the same issue locally on 2,008 `rule_73` configurations. It recovered 363 persistent stationary oscillators with periods `18..120` that the source detector had left bounded but unclassified. All source detections retained their original type and period, and the 165 period-12 cases used by Fases 81--87 were unchanged.

Fase 90 then reruns the two complete historical populations rather than extrapolating from that local cohort:

Cohort	Configurations	Confirmed misses
Background templates of length 1, 2, and 4	1,927,680	160
Primitive length-8 backgrounds	3,855,360	3,136
Total	5,783,040	3,296

Stage A reproduces 445,897 unique historical positive identities and applies the frozen counterfactual detector through period 120. The replay gate ignores only exact duplicate rows and order-dependent catalog annotations; physical fields remain part of the identity. It yields 3,296 bounded candidates above the historical cap. Stage B independently evolves exactly those 3,296 cases to `t=1000` and confirms every one with the same stationary type, minimal period, and zero drift. Candidate reconciliation finds no missing, extra, or duplicate physical identities and no type, period, or drift mismatches.

The recovered periods are:

Period	18	24	26	30	40	42	48	60	66	90	120
Cases	1,703	750	5	545	53	18	44	105	31	15	27

All confirmed misses occur in `rule_73` (1,623) or `rule_109` (1,673). They represent 0.0570% of the 5,783,040 configurations, so the absolute rate is small but the incompleteness is systematic rather than a local anomaly. The Fase 90 status is `GLOBAL_PERIOD_CAP_FALSE_NEGATIVES_CONFIRMED`.

This audit corrects the interpretation of the historical catalogs: a negative row under the source protocol means no oscillator was detected with period at most 16, not that no bounded periodic oscillator existed. It does not invalidate the previously detected oscillators or the $T=12$ causal analyses. Recurrence is computationally confirmed only through the frozen Stage-B horizon and period cap 120, and the result is exhaustive only for the two stated historical cohorts; it is not a theorem for arbitrary time, seed length, or background family.

7.40 Long-period attractor quotient and cut-coverage law (Fases 91--105)

The 3,296 confirmed long-period cases in Section 7.39 are input descriptors, not necessarily distinct attractors. Fases 91--93 separate three equivalence levels before interpreting basin occupancy. Joint background/defect cycles are canonicalized by a shared temporal rotation over the least common multiple of their periods; morphology, reflection, and `rule_73/rule_109` black/white conjugacy remain separate quotients rather than being folded into strict identity.

Representation level	Count	Interpretation
Confirmed input descriptors	3,296	Stage-B cases from Fase 90
Strict physical attractor classes	192	Joint background/defect periodic orbits
Defect-morphology classes	123	Translation-normalized defect cycles
Strict physical initial states	1,829	Absolute background plus absolute initial defect
Encoding aliases removed	1,467	Duplicate descriptions of the same initial state

The largest strict class contains 175 input descriptors before initial-state deduplication and 77 physical initial states afterward. Determinism is exact: no deduplicated initial state maps to more than one strict attractor class. Fase 92 then constructs the black/white conjugate trajectory for every case. All 3,296 defect trajectories, periods, drift values, and conjugacy-quotient hashes match exactly, closing all 123 quotient classes across both rules. Only 160 pairs are directly present in both frozen catalogs; 3,096 require the omitted complementary background phase, and 40 additionally require the zero-IC diagnostic excluded by the historical generator. The conjugacy result is therefore an exact deterministic symmetry check, while the original cross-rule frequency imbalance is a sampling asymmetry rather than independent statistical replication.

Fases 94--97 next construct a local intervention geometry. Each physical initial state is represented by its absolute eight-bit central word inside one of 48 fixed (`rule`, `background`) cubes. Edges are undirected Hamming-1 flips; they are interventions in initial-condition space, not time transitions of the CA. The 1,829 observed long-period states generate 14,632 one-bit interventions:

Hamming-1 outcome	Interventions
Same strict long-period class	1,818
A different long-period class	3,484
Explicitly represented outside the long-period set	9,249
Zero word excluded by the historical generator	81

Completing all 48 Q8 cubes from the frozen Stage-A ledgers distinguishes fragmentation within a cube from separation across backgrounds. Of the 192 strict attractor classes, 51 are connected within one cube, 21 are fragmented only across cubes, 41 are fragmented within at least one cube, and 79 combine both forms. Across 979 component pairs, the first connecting filtration level is the target class itself for 627 pairs, the addition of other long-period classes for 350, and the addition of represented short-period/static states for two.

The two final cases require the `SPAN_ESCAPE` category, but not a unique state: both have vertex and edge connectivity eight in the full non-zero Q8 graph. Thus category necessity and single-state bottlenecks are distinct.

Within the 627 first-level pairs, 219 targets have pair-specific vertex and edge connectivity one ($\kappa_v = \lambda_e = 1$); the other 408 are already redundantly connected. Adding the historical-source population rescues all 219 fragile targets, whereas the static-period-1 population rescues none. Period families show different raw unit-rescue rates, but they also expose different geometries. Fase 101 freezes a 13-variable, period-blind geometric signature and matches only within the same target. It covers 31,682 of 43,425 unit exposures (73.0%). All 4,090 matched geometry strata are outcome-homogeneous, and all 7,281 within-stratum period comparisons are exact ties for both vertex and edge rescue. The remaining 11,743 unmatched exposures are retained as an explicit coverage limit rather than extrapolated.

Fase 102 replaces that matched association with an exact topological law on the full 43,425-exposure population. For a fragile baseline graph G , let C_V and C_E be the complete sets of pair-specific single-vertex and single-edge cuts between its terminal components. For an added historical state x , let $b_V(x)$ and $b_E(x)$ count the cuts bypassed by the new Hamming-1 edges. Then, under the frozen Q8 protocol,

$$\begin{aligned} \kappa_v(G + x) \geq 2 & \text{ iff } b_V(x) = |C_V| \\ \lambda_e(G + x) \geq 2 & \text{ iff } b_E(x) = |C_E|. \end{aligned}$$

Metric	True positive	True negative	False positive	False negative
Vertex redundancy (κ_v)	1,505	41,920	0	0
Edge redundancy (λ_e)	1,566	41,859	0	0

The equivalence is not inferred from correlation: the critical cuts are re-enumerated from the frozen graph, every intervention is checked by direct cut removal, and a separate max-flow computation reproduces all 43,425 outcomes. The mechanism is exact because adding a node and edges cannot create a new separator whose removal disconnects a baseline graph that was already connected; rescue therefore occurs precisely when all original unit cuts are bypassed.

Fase 103 tests the first collective extension without changing that graph definition. It enumerates all 404,054 unordered pairs in 142 strata where a complete historical period family rescues but no singleton does for at least one metric. Structural cut coverage and fresh max-flow agree on every pair.

Metric	Collective strata	Exactly cardinality 2	Minimum only bounded >=3	Rescuing pairs
Vertex redundancy	126	69	57	454
Edge redundancy	139	68	71	470

The mutual Hamming-1 edge between the two added words is necessary in 83 vertex rescues and 86 edge rescues. Exactly 2 is a proof of minimum cardinality because all constituent singletons were audited first.

Fases 104--105 complete the unresolved ladder without changing targets, terminal components, graph filtration, or rescue criterion. Fase 104 enumerates all 3,061,466 unordered triples in the 73 strata left unresolved by pairs. Direct cut removal and fresh max-flow agree on every trial:

Metric	Strata entering triple audit	Exactly cardinality 3	Remaining after triples	Rescuing triples
Vertex redundancy	57	41	16	180
Edge redundancy	71	40	31	192

Fase 105 then enumerates all 24,362,850 unordered quadruples in the remaining 32-stratum union, comprising 20,638,850 vertex and 19,941,575 edge trials. It finds at least one rescuing quadruple in every unresolved metric-specific stratum:

Metric	Strata entering quadruple audit	Exactly cardinality 4	Still unresolved ≥ 5
Vertex redundancy	16	16	0
Edge redundancy	31	31	0

All 77 vertex-rescuing and 103 edge-rescuing quadruples require at least one internal Hamming-1 edge among the four added words. Across the full audit, Route A (exhaustive cut removal) and Route B (independent integer max-flow) have zero disagreements, all reserved and out-of-scope ledger bits remain zero, and independent decoding reproduces every aggregate. Combining the pair, triple, and quadruple stages gives the complete minimum-cardinality partition for the collective-only strata: vertex 69/41/16 and edge 68/40/31 at cardinalities 2/3/4, respectively. The combined status is QUADRUPLE_SYNERGY_ATLAS_BUILT.

These results are mechanistic but protocol-bounded. The exact law covers 219 fragile targets inside 48 frozen length-8 Q8 intervention cubes and the historical state populations represented there. It is not a theorem about all ECA basins, arbitrary initial-condition widths, temporal causality, or interventions outside the frozen catalogs. Exact cardinalities two through four refer only to the frozen collective strata and rescue definition; they do not establish a universal upper bound for other ECA rules, word lengths, target classes, or intervention families.

7.41 Minimal-rescue motifs and exact QUBO compilation (Fases 106--107)

Fase 106 freezes geometry before joining outcomes and classifies every candidate intervention set from the pair, triple, and quadruple ledgers by the exact isomorphism class of its induced Hamming-1 graph. Two independent classifiers -- one based on graph invariants and one based on canonical adjacency under all vertex permutations -- agree on all 27,828,370 records. The closed triangle-free motif vocabulary contains two classes at cardinality two, three at cardinality three, and seven at cardinality four.

Cardinality	Candidate sets	Motif classes	Vertex rescues	Edge rescues	Internal-edge dependent (vertex / edge)
2	404,054	2	454	470	83 / 86
3	3,061,466	3	180	192	98 / 98
4	24,362,850	7	77	103	77 / 103

The resulting 1,476 metric-specific minimal rescues retain three distinct properties rather than collapsing them into a single label: complete coverage of the frozen critical cuts, full rescue under direct cut removal and fresh max-flow, and dependence on internal Hamming-1 edges among the added states. Every rescue covers the corresponding cuts and reproduces the verified graph outcome. At cardinality two, 371 vertex and 384 edge rescues remain after all internal edges are removed; at cardinality three the corresponding counts are 82 and 94. No cardinality-four rescue survives removal of all internal edges. Motif classes remain heterogeneous where the frozen data make them so; the atlas does not reinterpret a frequent motif as a causal explanation.

Fase 107 compiles the verified rescue hypergraphs into independent sparse quadratic unconstrained binary optimization (QUBO) models. For one frozen target-period-metric instance, let x_i select historical state i , let z_h select one verified minimal rescue hyperedge h , and let every hyperedge have the already proved cardinality d . With unit node costs and penalty $P=d+1$, the compiled objective is

$$E(x, z) = \sum_i x_i + P (1 - \sum_h z_h)^2 + P \sum_h z_h (d - \sum_{\{i \text{ in } h\}} x_i).$$

Using binary identities, this produces constant P , linear x_i coefficients $+1$, linear z_h coefficients $P(d-1)$, pairwise $z_h z_{h'}$ coefficients $+2P$, and incidence coefficients $z_h x_i = -P$, with no $x_i x_j$ terms. For any selected hyperedge, each included x_i has net coefficient $1-P < 0$ and each excluded x_i has coefficient $+1$; therefore the unique minimum for that z_h is exactly its incidence vector, with energy d . The one-hot penalty makes $z=0$ cost $P > d$ and makes two or more active hyperedges non-optimal. Thus the ground-state degeneracy is exactly the number of verified rescue hyperedges, not an empirical solver observation.

QUBO compilation quantity	Value
Independent target-period-metric models	265
Accumulated x variables	17,624
Accumulated z variables	1,476
Accumulated variables	19,100
Variables per model	9--172
Sparse nonzero terms	32,861
Certified ground states	1,476

The 265 models partition as 69/41/16 vertex-connectivity instances and 68/40/31 edge-connectivity instances at cardinalities two, three, and four. Two independent integer-energy routes reproduce every coefficient and every certified ground state with zero discrepancies. This compilation does not run cellular-automaton simulations, discover new rescues, execute a quantum or annealing device, or establish a speedup. It supplies an exact, reusable QUBO representation of the already verified finite rescue problem under unit node costs; weighted costs, larger intervention families, and hardware behavior remain outside scope.

7.42 Mechanism-conditioned allocation of a fixed Hamming-incidence budget (Fases 108--112)

Fase 108 conditions the 1,476 verified minimal rescues on both their frozen Hamming-1 motif and whether rescue survives removal of every internal edge of the intervention set. The binary mechanism label is therefore an exact audit of internal-edge dependence, not a latent class inferred from geometry. Motifs without internal edges are logically forced to the external-attachment label; three cardinality-four motifs have no observed minimal rescue; and only K_2 and K_2+I contain both labels in the frozen census.

Observed motif	Rescues	External attachment	Internal-edge dependent	Classification
2I	701	701	0	forced
K_2	223	54	169	variable
3I	154	154	0	forced
K_2+I	96	22	74	variable
P_3	122	0	122	complete separation
K_2+2I	35	0	35	complete separation
$2K_2$	48	0	48	complete separation
P_3+I	67	0	67	complete separation
P_4	30	0	30	complete separation

The clean within-instance comparison is the K_2 subset. It contains 223 rescues in 101 target-period-metric instances, of which 24 instances contain both mechanisms (122 rescues). For a rescued pair $H=\{u, v\}$, let v_i be the period-specific candidate universe reconstructed from the certified pair ledger and let $G_{\min}$ be

the frozen bridge graph reconstructed from the earliest allowed physical level. In all 24 mixed instances, H is a subset of V_i and V_i is disjoint from $G_{\min}$.

Each endpoint of a K_2 pair has seven Hamming-1 incidences other than the internal edge. Define A_V as the total incidences from the two endpoints to V_i outside H , A_G as the incidences to $G_{\min}$, and A_R as the unassigned residual. The disjointness gives the exact identity

$$A_V(H) + A_G(H) + A_R(H) = 14.$$

The comparison gives every mixed instance equal weight. For each feature, Δ_i is the mean among external-attachment rescues minus the mean among internal-edge-dependent rescues in the same instance.

Allocation feature	Mean Δ_i	Median Δ_i	Positive / zero / negative instances
A_V (candidate universe)	-1.220833	-0.500000	2 / 0 / 22
A_G (bridge graph)	1.816667	1.550000	24 / 0 / 0
A_R (residual)	-0.595833	-0.700000	2 / 8 / 14

Thus the conditioned difference is an allocation shift, not an increase in the total budget: external-attachment rescues allocate fewer incidences toward V_i and more toward $G_{\min}$ on average. After centering both variables within instance and giving each instance equal total weight, the descriptive Pearson correlation between A_V and A_G is -0.824731 . This correlation is partly mechanically constrained by the fixed budget and is not evidence of statistical independence, dependence, or causality.

Fase 111 tests whether A_G alone can classify individual rescues without reusing the held-out instance to choose its threshold. The apparent aggregate performance is dominated by 77 monolabel instances. In the 24 mixed instances, where the two mechanisms must actually be distinguished, the weighted balanced accuracy is only 0.647500 and internal sensitivity is 0.295000 . Moreover, A_G itself was selected after inspecting this census, so this internal leave-one-instance-out analysis is neither external nor prospective validation. The geometric allocation should therefore be read as a conditioned census result, not as a predictive rule.

Fase 112 records a separate post-selection calibration of rescue multiplicity. Because its multiplicity threshold, statistic, and tail direction were chosen after observing the same outcomes, it is retained as an auditable exploratory artifact and is not used as a confirmatory paper claim or a formal rejection of an exchangeable null model.

8. Observer Artifacts and Pipeline Equivariance

The ZUSE pipeline contains two classes of observer artifact that the atlas identifies and characterizes. Framing these artifacts as results -- not just implementation limitations -- is important: they define the boundary between what the system measures reliably and what it does not.

8.1 Mirror asymmetry in `tipo_unico` (Fase 6b)

`rule_110` and `rule_124` are left-right mirrors of each other: the rule table for `rule_124` is obtained by reflecting every neighborhood $f(l,c,r) \rightarrow f(r,c,l)$ in the `rule_110` table. Under periodic boundary conditions, the two rules produce physically equivalent dynamics up to spatial reflection.

Despite this equivalence, `tipo_unico` can fire asymmetrically: it may be accepted for one orientation and rejected for the other, depending on which structure types the heuristic observers label from the specific frame sequence. Since `tipo_unico` counts whether exactly one structure type appears, its value depends on the labeling convention of the observers, not only on the CA dynamics.

`tipo_unico` is retained in the atlas for its exploratory value: it reliably distinguishes runs with homogeneous structure populations from runs with mixed populations. It should not be used as evidence of physical left-right asymmetry.

8.2 Translation non-equivariance of the dedup pipeline (Fase 19)

ECA with periodic boundary conditions is translation-invariant: shifting the initial condition by any number of cells produces the same dynamics up to a spatial shift. A pipeline that correctly identifies physical structures should therefore return the same structure count for all translations of the same IC.

Fase 19 tested this directly: 64 single-bit ICs for `rule_54`, one active bit at each position `k = 0..63`, with `width = 64` and `steps = 96`. The ECA frames were confirmed translation-invariant (frame identity after shift normalization: `True`). The observer/dedup pipeline was not:

metric	range across 64 positions
<code>dedup_structure_count</code>	15..24 (29 distinct result classes)
<code>raw_structure_count</code>	45..72
law signature	<code>temporal_scale_stability</code> (all 64 identical)
<code>analysis_status</code>	ok (all 64)

metric	value
ICs tested	64 single-bit positions (<code>k = 0..63</code>)
ECA translation-invariant	True
Unique result classes (observer)	29
dedup range observed	15-24
Noise gate threshold	> 40
Closest single-bit IC to gate	<code>k = 55</code> , <code>dedup = 24</code>
Complex-IC reference (Fase 13)	<code>dedup = 32..39</code>
Law signature (all 64 positions)	<code>temporal_scale_stability</code>
Conclusion	noise-gate crossing requires complex IC geometry; single-bit ICs stay at least 16 below gate

The mechanism is a boundary interaction: `rule_54` produces wide-spreading patterns that cross the periodic frame boundary. The dedup algorithm's handling of cyclic-span structures varies depending on the absolute IC position relative to where structure boundaries fall on the lattice. The result is a position-dependent count that is not a translation-equivariant physical observable.

Fase 25 isolates a second instance of the same geometric failure mode in the local-oscillator detector. Co-translated background and XOR perturbation orbits are identical in 80/80 runs, but `linear_shape` loses 22 moving signatures when the localized difference straddles positions 255 and 0. Circular canonicalization restores 80/80 signatures. This directly attributes the non-equivariance to linear treatment of a periodic lattice.

8.3 Implications for the atlas

Both artifacts are bounded in their effect:

- `tipo_unico` asymmetry is a labeling artifact, not a count artifact. It affects which laws are accepted, but only for runs where the structure population is near the one-type boundary. Runs with clearly homogeneous or clearly mixed populations are unaffected.
- Dedup non-equivariance affects absolute structure counts but not law signatures. In Fase 19, all 64 translated ICs produce identical law signatures despite varying dedup counts. The noise gate (`dedup >`

40) is never approached by single-bit ICs, and law evaluation depends on count magnitude only through the gate.

The atlas therefore relies on law signatures as the primary evidence unit. Absolute dedup counts appear in world profiles as context and should be interpreted with the translation-equivariance caveat. Future work on symmetry-invariant observers would remove both artifacts.

9. Limitations

9.1 Fixed protocol parameters

The atlas is valid for the parameter regime used: `width = 64` (formal profiles), `width = 128` (rule_108 oscillator), `steps` roughly `24..200`, and the IC protocols defined per world. The two calibrated thresholds -- `frontera_temporal` upper bound `0.4352` and `temporal_scale_stability` threshold `19.03` -- were fit on data from this regime. Applying the atlas to significantly different widths or step counts requires recalibration. This is not a flaw in the methodology; it is the expected scope of an empirically grounded atlas.

9.2 Heuristic observers

The observer stack uses geometric heuristics to label structures as `glider`, `bloque`, or `oscilador`. These heuristics are not derived from first principles and are not provably complete or sound for arbitrary ECA dynamics. As shown in Section 8, they are not translation-equivariant for wide-spreading patterns and are not mirror-invariant for `tipo_unico`. The atlas is built on law signatures, which are more robust than absolute observer counts, but the underlying observers remain heuristic. Replacing them with symmetry-invariant observers would be a meaningful improvement.

9.3 Bounded local oscillator protocol

The uniqueness claim for `rule_108` holds under a specific protocol: quiescent zero background, stationary exact periodicity (no drift), IC words of binary length `1..12` (502 words of length `1..8` plus 7,676 words of length `9..12`), and period detection window `2..16` with local span `<= 32`. IC words longer than 12 or longer detection periods are outside the current stationary protocol. A separate periodic-background sweep (Section 7.6) confirms that non-zero backgrounds substantially change the oscillator landscape; those results define a different regime and do not apply to the zero-background uniqueness claim. Moving oscillators under quiescent zero background were searched in a companion sweep -- Section 7.5 reports the result. The claim is therefore: no other quiescent ECA rule produces a stationary local-period oscillator under this protocol. It is not a claim about ECA oscillators in general.

The period-8 background sweep (Section 7.7) is also bounded by background phase. Fase 25 completed the strict co-translation test: the physical orbits are exact translations in 80/80 runs, while the original linear observer recovers 58/80 signatures because of cyclic boundary crossings. Circular canonicalization recovers 80/80. Fixed-IC phase dependence remains in all 10 sampled cases after correction, so it is physical alignment sensitivity rather than pure observer non-equivariance.

Fase 26 strengthens the longest-period result without converting it into a general theorem. All 20 minimal `T=15` rule/background representatives persist through step 900 and share a background temporal period of 3, but only 23/160 background phases and 4/134 one-bit IC mutations retain `T=15`. The result is therefore robust in time and narrow in basin under the tested protocol.

Fase 27 establishes the state-cycle mechanism of the 5:1 locking computationally. After burn-in, the localized defect $D(t) = X(t) \text{ XOR } B(t)$ cycles through exactly five distinct states under F^3 across all 20 minimal representatives (Section 7.9). The 5:1 ratio is therefore the cycle length of the defect under the three-step

operator, not a spatial resonance. The symbolic derivation of why this five-cycle appears in `rule_73/rule_109` but not in other rules over `T_bg=3` backgrounds remains open.

Fase 28 proves the black/white conjugation of the defect dynamics analytically and verifies it exhaustively at local and orbit levels (Section 7.10). It also rules out one tempting simplification: no fixed sparse subset of induced local transitions is shared by all 100 `F^3` edges. The remaining symbolic problem requires phase-aware or higher-order block variables.

Fase 29 and Fase 30 sharpen that boundary. The `T=15` cycle is not a pure defect-only dynamic: canonical defect shapes are background-specific already at `w=0`, and no nontrivial fixed local block signature is shared across all backgrounds. But the variation is not arbitrary. The 20 representatives collapse into 13 phase-rotated defect-cycle families, including two families shared across `rule_73` and `rule_109`. The open problem is therefore not simply "find the local rule"; it is to map the temporal background orbit and IC alignment to a finite shape family and phase offset.

Fase 31 and Fase 32 provide the current compact description of that map. No tested background-only descriptor shorter than the full temporal orbit predicts the 13 families globally. However, conditioned on rule identity, the circular multiset of length-4 background subwords separates the families, and rotational variants preserve the predicted family in 140/140 co-translated runs. The same rotations fail almost completely with fixed IC placement (3/140 detections), so IC/background alignment is an essential physical variable, not a bookkeeping detail. Fase 33 then shows that length 8 contains no unseen same-descriptor backgrounds outside the known representatives and their rotations; the two length-4 descriptor collisions that do exist are already family-preserving. Fase 34 moves to primitive length-9/10 backgrounds with `T_bg=3` and finds 90 `T=15` detections across 8 external backgrounds. This confirms that the five-to-one mechanism generalizes beyond length 8, while leaving the variable length descriptor problem open.

Fases 35--38 replace visual shape families with explicit macro-transition tables and effective orbit embeddings. The induced `F^3` transition table is a sufficient discriminator and refines the 13 visual families. The `F00` table identity is explained by convergence to the same effective period-3 background orbit, but that explanation is too coarse globally: each rule has only one canonical period-3 orbit across the confirmed set. The first sufficient post-burn-in descriptor is therefore `(rule, sample_orbit_step, sample_rotation_offset, defect_state0)`, where `defect_state0` is the first sampled stable-cycle defect state.

Fase 39 tests whether that final measured variable can be predicted compactly from pre-burn-in data. The answer is negative under the tested descriptors. All 20 representatives enter the stable five-cycle by `t=12` and 15/20 enter at `t=3`, but exact pre-burn-in predictors either rely on post-hoc entry time or collapse into singleton case identifiers. The remaining formal problem is therefore precise: predict `defect_state0` from the raw background/IC pair without effectively replaying the early transient.

Fase 40 shows that "effectively replaying" does not require the full system. A strict 25-cell causal cone centered on the IC and simulated for 12 steps matches the full defect state at `t=12` in 20/20 representatives and, after phase projection, recovers `defect_state0` at `t=81` in 20/20. The compression relative to the full 256-by-81 simulation is 69.1x. Thus the unresolved pre-burn-in component is locally causal rather than global: the open question is how to replace the 25-cell, 12-step cone computation with a symbolic rule.

Fase 41 audits whether that local computation contains an obvious smaller support. It does not. All 25 initial cone inputs are required for the final active defect support, all eight ordinary ECA table entries are used in every representative, and the induced background/defect tables are dense (49..62 of 64 possible keys). The cone has only structural pruning for active outputs (234..310 internal nodes instead of 325), so the remaining symbolic problem is Boolean minimization of a dense local circuit, not discovery of a sparse subtable or smaller causal support.

Fase 42 verifies the input-support conclusion at the Boolean-function level. Natural-order ROBDDs for the active localized output functions contain all 25 support variables in every representative, with 17,141..36,966 reachable nodes for active outputs and 51,539..53,901 for the full 25-bit vector. Variable support is semantic, not merely an artifact of the natural BDD order, so this rules out Boolean input elimination. It does not claim a globally minimal BDD size over all variable orders.

Fase 43 then tests the most direct BDD-size escape route. Existing `natural`, `reverse`, and `center_out` orders show only weak useful sensitivity: `reverse` improves total active-output nodes by 0.5%, while `center_out` is worst in 20/20 representatives. A checkpointed one-pass SIFT search on the most favorable representative evaluates 580 orders and improves 16,061 active-output nodes to 16,056, far above the 10,000-node compression gate. This does not prove global BDD-size optimality, but it rules out simple variable-ordering as the missing symbolic shortcut.

Fases 44-45 then switch representation class from BDDs to ANF polynomials. They show that the Boolean cone is not algebraically uniform. Active-output degree obeys $\text{degree} = 24 - |\text{rel_pos}| + \text{epsilon}$, with `epsilon` in $\{0,1\}$ and zero exceptions over 174 outputs; monomial counts decay with $\log_{10}(\text{monomials}) \approx 7.241925 - 0.307283 * |\text{rel_pos}|$ ($R^2=0.998197$). This spatial gradient is centered on the defect and is orthogonal to the BDD negative results. It opens a sharper symbolic question: whether the gradient can be derived from cone geometry and the background orbit, and what determines the residual epsilon for distances ≥ 2 .

Fase 46 tests that residual directly. After removing `dist=0` and `dist=1`, where `epsilon=0` trivially, it evaluates 141 active outputs using static features from rule identity, signed position, local background bits, family id, and defect phase. The best single feature reaches only 64.89% leave-one-representative-out accuracy, and a depth-3 decision tree falls from 73.05% training accuracy to 55.65% leave-one-representative-out accuracy. Thus the epsilon bit remains residual under the current static feature class.

Fase 47 then changes the feature class from static descriptors to dynamic ANF growth profiles. Recomputing ANF degree and monomial count at each cone layer `t=1..12` yields zero mismatches against the final Fase 44 degrees at `t=12`. The single feature `degree_growth_slope` predicts `epsilon` with 94.90% leave-one-representative-out accuracy. The result should not be read as a pre-computation shortcut, because the feature uses the full temporal trajectory through the final layer; it is instead a dynamic law explaining where the one-bit residual lives.

Fase 48 then asks whether that dynamic law appears before the final cone layer. The answer is negative under the tested horizon protocol: the future-blind `degree_growth_slope_K` feature reaches 79.47% LORO accuracy at $K=11$, then jumps to 94.90% at $K=12$. Thus the epsilon bit is a full-horizon ANF-growth property rather than an early dynamic shortcut.

The external ANF holdout in Section 7.38 is likewise protocol-bounded. It adds 25 `rule_73` cases on 18 primitive length-8 backgrounds, but only nine natural-period witnesses, all at `T_local=12`. Three retain the comparison signature at one neighboring horizon, while six do not and none survives at horizon 16. The fact that 0/27 control measurements become positive reduces the likelihood of a free threshold effect, but the sample still covers only one non-`rule_109` rule and one background length. It therefore falsifies universal necessity of center mediation while supporting only partial, background-conditioned generalization of the gradient.

Fases 89--90 quantify a separate detector-completeness limit. The historical periodic-background detector searched exact periods only in `2..16`; a full re-sweep of its 5,783,040 frozen configurations recovers 3,296 stationary oscillators with periods 18..120. All occur in `rule_73` or `rule_109`, and all are confirmed through `t=1000`. This correction does not change any historical positive classification or the 165-case `T=12` cohort used by Fases 81--87, but historical negative rows must be interpreted as detector-bounded. The extended check remains finite in horizon and period and does not establish absence beyond period 120.

Fases 91--112 add a topological intervention layer, but its exactness must not be confused with universal scope. The cut-coverage equivalence is proved computationally for 43,425 unit additions to 219 fragile targets inside 48 frozen length-8 Q8 cubes. These graphs encode Hamming-1 changes of the central eight-bit initial word; they are not CA time-transition graphs. The 31,682 geometry-matched exposures cover 73.0% of the unit population, so the period-conditioned tie in Fase 101 is not extrapolated to the unmatched 11,743. Fase 102 removes that matching restriction only for the cut-coverage law, not for arbitrary observables or other intervention spaces. Fases 103--105 prove minimum rescue cardinalities two, three, or four by exhaustive enumeration in the collective-only strata, with every lower cardinality already audited. This closes the ladder only for the frozen state populations and graph definition; it is not an upper bound for arbitrary intervention spaces. Fase 106 classifies the induced Hamming-1 motifs of those same frozen interventions; the motifs are descriptive graph classes and are not promoted to universal causal primitives. Fase 107 compiles the verified minimal rescues into unit-cost QUBOs. The compilation is exact for its 265 finite instances, but no quantum hardware or annealing solver was run, and no quantum speedup, practical advantage, or transfer beyond the frozen intervention family is claimed. Fases 108--110 further condition those same rescues by mechanism and geometry; they do not add a new population. Their fixed-budget comparison is restricted to 122 rescues in 24 mixed `k2` instances. Fase 111 is an internal diagnostic whose feature was selected on the same census, and Fase 112 is explicitly post-selection. Neither supplies external validation, a population generalization, a causal effect, or a confirmatory significance test.

9.4 Empirical atlas, not axiomatic classification

The world categories are induced from observed law signatures across a finite number of seeds and step counts. A world classified as `multiregimen-productivo` on 6..15 visits could exhibit different behavior at larger scale, with different IC distributions, or under longer runs. The categories are stable empirical summaries, not theorems. `rule_90` is a clear example: it is classified as `multiregimen-escala-dependiente` because high-scale visits become silent under the current protocol, but the underlying XOR dynamics have algebraic structure that the current seven laws do not capture.

Fase 20 gives the same warning for `frontera_temporal`: 24 additional rules were rich in `frontera_temporal` at sweep scale, but the top four failed long-journal validation. Category assignment is therefore protocol-scale dependent; short-scale richness is candidate evidence, not atlas-grade classification.

Fase 21 gives an analogous warning for `periodicidad`. With random ICs, production `periodicidad` is rare in ECA; with explicitly periodic IC families, it appears in 207/256 rules. The law is therefore not inaccessible to ECA, but it is strongly conditioned by the IC family. Atlas rows should be read as claims about a stated rule/IC/scale protocol, not as intrinsic properties of the rule table alone.

9.5 PySR symbolic regression -- negative result

The decision-tree analyses (Section 4, temporal calibration) provide strong empirical signal but not closed-form symbolic expressions. PySR was planned as a follow-up to produce interpretable formulas for the calibrated thresholds and fragility spectra. PySR 1.5.10 was subsequently unblocked. A full regression over 15 atlas worlds (five features: `mean_laws`, `peak_diversity`, `noise_ratio`, `non_empty_ratio`, `f_core`; target `f_total`; 40 iterations) produced a best expression of complexity 9 with `MSE = 0.035`, above the threshold for a paper-worthy finding. The dominant predictor is `f_core`; the residual is driven by `rule_108` as a structural outlier (`f_gap = 0.945`) that is not predictable from aggregate features without a mechanism feature. The symbolic layer remains incomplete, and this negative result is consistent with the `f_core/f_gap` separation documented in Section 6.

10. Next Work

10.1 Symmetry-invariant observers

The two observer artifacts identified in Section 8 -- `tipo_unico` mirror asymmetry and dedup translation non-equivariance -- share a root cause: the heuristic observers do not encode the symmetries of the underlying CA. A natural next step is to build observers that canonicalize structure representations under spatial reflection and translation before counting. This would make `tipo_unico` a mirror-invariant physical property and dedup counts stable across IC positions, strengthening the evidential basis for both the atlas and the fragility measurements.

10.2 Extended local oscillator search

The `rule_108` uniqueness result holds under the current stationary protocol (zero background, exact period, IC words of length ≤ 12 , span ≤ 32). Several controlled extensions have now been completed:

- **Moving oscillators:** completed. A companion sweep over all 128 quiescent rules found eight rules producing minimal period-2 speed-1 gliders (Section 7.5). No longer-period or slower-speed moving oscillators were found under this protocol. The follow-up length-9..12 sweep found no additional stationary or moving oscillator rules.
- **Longer IC words beyond 12:** extend the IC sweep from length 12 to length 16 or beyond to test whether substantially wider seed patterns produce oscillators in rules that failed the length-12 protocol.
- **Non-zero backgrounds:** completed. A sweep over 15 non-zero periodic backgrounds (template lengths 1, 2, 4) across all 256 rules and 502 IC words (1,927,680 runs) found that the periodic-background regime is substantially richer than the quiescent regime: 30 stationary rules, 36 moving rules, including period-4 oscillators and speed-0.5 gliders not present under zero background. `rule_108` persists under all-one background with the same motif. Full results in Section 7.6.
- **Period-8 backgrounds:** completed. A sweep over 30 primitive length-8 binary necklaces (3,855,360 runs) found 4 new stationary rules, 19 new moving rules, five new period classes ($T=6, 8, 10, 12, 15$), and a new glider speed ($2/3$ cell/step, $T=3$). All 10 sampled results were background-phase dependent after circular-geometry correction. Fase 25 confirms exact co-translation equivariance in 80/80 runs with circular shape canonicalization. Full results are in Section 7.7.
- **$T=15$ anatomy:** completed. All 221 detections belong to the reflection-symmetric, black/white-conjugate pair `rule_73/rule_109`. The 20 minimal rule/background representatives remain exact through step 900 and lock at five times the background temporal period ($15/3$). The basin is narrow: 23/160 background phases and 4/134 one-bit IC mutations retain $T=15$. Full results are in Section 7.8.
- **Five-state locking mechanism:** completed. The localized defect cycles through five distinct states under F^3 . The ratio $T_{\text{local}}/T_{\text{bg}}=5$ is the defect cycle length, confirmed in all 20 minimal representatives. Full results are in Section 7.9.
- **Induced defect algebra:** partial. The exact law $\text{delta_f}(b,d)=f(b \text{ XOR } d) \text{ XOR } f(b)$ proves that the `rule_73` and `rule_109` defect orbits are identical under simultaneous black/white conjugation. Profiling all 100 F^3 edges rejects a universal sparse-entry explanation and points instead to spatial phase or higher-order block states. Full results are in Section 7.10.
- **Block locality and shape families:** partial-positive. Fixed local defect blocks do not explain the $T=15$ cycle: the active defect already carries background-specific context at $w=0$, and no nontrivial shared block signature exists across all backgrounds. The variation is nevertheless finite: 20 representatives collapse into 13 phase-rotated shape families, with two families shared across the conjugate rules. Full results are in Section 7.11.
- **Compact $T=15$ state variable:** completed for the known representative set. No compact background-only descriptor predicts all 13 families globally, but the rule-conditioned descriptor `subpatterns_1en4` separates

- all backgrounds per rule. Rotation generalization validates the compact state variable (`rule`, `subpatterns_len4`, `IC/background alignment`): 140/140 co-translated rotations preserve both $T=15$ and the predicted family, while fixed-IC rotations collapse to 3/140 detections. Full results are in Section 7.12.
- **External $T=15$ validation:** completed for targeted primitive length-9/10 backgrounds. Fase 33 proves that length 8 has no unseen same-descriptor backgrounds outside rotations of the known set. Fase 34 then targets the 66 primitive length-9/10 backgrounds with $T_{bg}=3$ under `rule_73/rule_109` and finds 90 $T=15$ detections across 8 external backgrounds. Full results are in Section 7.13.
 - **Transition-table and embedding analysis:** completed for the 20 minimal $T=15$ representatives. The explicit F^3 transition table is a sufficient discriminator and refines the visual family partition. Family `F00` is explained by convergence to a shared effective period-3 background orbit, but canonical orbit identity alone is too coarse globally. The first sufficient post-burn-in descriptor is (`rule`, `sample_orbit_step`, `sample_rotation_offset`, `defect_state0`). Full results are in Section 7.14.
 - **Pre-burn-in entry phase:** negative delimiter. The defect enters the stable five-cycle quickly (all 20 by $t=12$, 15/20 at $t=3$), but the tested pre-burn-in descriptors do not give a compact predictor of the entry phase: exact descriptors are mostly singleton case identifiers or post-hoc measurements. Full results are in Section 7.15.
 - **Early causal-cone predictor:** positive. The strict $2t+1$ causal cone centered on the IC at $t=12$ (25 cells for 12 steps) matches the full defect state in 20/20 representatives and recovers `defect_state0` after phase projection in 20/20. This is a 69.1x compression relative to the full 256-by-81 simulation. Full results are in Section 7.16.
 - **Minimal cone-table audit:** negative/structural. The Fase 40 cone has no sparse table shortcut: induced `(b,d)->d_next` tables use 49..62 of 64 keys, all eight ordinary ECA entries are used, and all 25 initial cone inputs are required for the active localized output. The only reduction is structural: active-output dependency uses 234..310 internal nodes instead of the full 325. Full results are in Section 7.17.
 - **ROBDD cone audit:** negative/semantic. Reduced ordered BDDs under the natural cone order confirm that the active-output functions depend on all 25 inputs in every representative. Active-output BDDs have 17,141..36,966 reachable nodes; the full 25-bit vector has 51,539..53,901. This rules out Boolean input elimination, while not claiming globally minimal BDD size over all variable orders. Full results are in Section 7.18.
 - **ROBDD order search:** negative/bounded. The `reverse` order improves total active-output nodes by only 0.5% over `natural`, while `center_out` is worst in 20/20 representatives. A targeted 580-order SIFT pass on the best known representative improves 16,061 active nodes to 16,056 (0.031%), missing the 10,000-node compression gate. Full results are in Section 7.19.
 - **ANF gradient:** positive/structural. Exact ANF analysis of 174 active outputs shows that $\text{degree} = 24 - |\text{rel_pos}| + \text{epsilon}$, with `epsilon` in $\{0,1\}$ and zero exceptions, while monomial counts decay almost by a factor of two per cell from the cone center ($R^2=0.998197$). The epsilon residual is tested directly in Fase 46: after excluding trivial `dist=0,1` rows, the best single-feature predictor reaches 64.89% leave-one-representative-out accuracy, and a depth-3 tree reaches only 55.65% under the same validation. The residual remains unexplained by static rule/position/background/family features. Fase 47 resolves the residual at the dynamic-profile level: `degree_growth_slope` over $t=1..12$ predicts epsilon with 94.90% leave-one-representative-out accuracy, with zero mismatches against Fase 44 at $t=12$. Fase 48 shows that this is a full-horizon effect: $K=11$ reaches only 79.47% LORO accuracy, while $K=12$ jumps to 94.90%. Fase 49 validates the gradient on external length-9/10 $T=15$ witnesses (0/63 exceptions, $R^2=0.998263$). Fases 50--53 then test specificity outside the original $T=15$ set: compact $T=2$ baselines lack enough active support; `rule_109/T=10` reproduces the $T=15$ slope within 0.13%; and the external `rule_54`, `rule_94`, and `rule_133` families are flat at their natural periods. Fase 54 then tests additional `rule_73/rule_109` family witnesses and temporarily isolates the known `rule_109/background 1011/T=10` case. Fase 55 replaces that local conclusion with a full catalog census: 66 non- $T15$ groups with `span >= 11`, 128 ANF

measurements, two new `rule_109/T=12` natural-period witnesses, and two new `rule_109/T=8` acceptable-horizon witnesses. No `rule_73` or external-rule case becomes a strong/acceptable witness. Fase 56 audits the structural cause: orbit symmetry is partial, but `rule_109` has a center-mediated local ANF (`LC`, `CR`, `LCR`) whereas `rule_73` has an isolated center term and an `LR` term without the center. Fase 57 then tests this candidate against the full catalog: center mediation is necessary for the 5 positive witnesses (0/5 positives have `center_mediated=False`), but not sufficient, because 12 center-mediated `rule_109` catalog cases are non-positive. Fase 58 tests the next candidate condition, period/horizon, inside `rule_109`: `T_local >= 8` captures all positives but produces four false positives, while `T_local=12` has no false positives but captures only 2/5 positives. The residual `bg=0011/T=8` shows that period/horizon is only a partial discriminator and points to background phase, IC placement, or alignment. Fases 59--60 then test that alignment layer: IC placement separates the three `T=8` cases, but the rule does not generalize to all 17 `rule_109` cases, and the only perfect separator is the lookup-like exact `ic_active_bits` pattern. Fases 61--62 then test dynamic alignment: no single descriptor separates all positives, but the minimal union `size_growth_total <= -3 OR center_drift_abs <= 0.0` captures 4/5 positives with zero false positives, leaving `bg=1100/T=8/word=00000110` as the universal residual. Fase 63 then tests four minimal monomial-level ANF interventions around `rule_109`; all four edits are algebraically valid, but none preserves a stationary oscillator on `bg=1100`, so the direct intervention test is blocked before cone-ANF measurement. Fase 64 then tests the Hamming-1 truth-table neighbourhood: seven neighbours have no oscillator support on `bg=1100`, while the only survivor, `rule_108`, collapses to a compact `T=2` mechanism with non-comparable ANF slope (-0.017102 , $R^2=0.264706$). Fases 65--66 then test the residual directly: aggregate context frequencies do not discriminate it from nearby controls, and a long-horizon audit confirms that it is a genuine persistent period-8 oscillator rather than a transient. Fases 77--79 then add a predeclared external-background holdout. Four new period-3 rules produce no comparable fits, but 9/25 `rule_73` cases on primitive length-8 backgrounds reproduce the gradient at their natural period `T=12`. Three of those nine survive at one neighboring horizon (`h=10` or `h=14`), while 0/27 neighboring-horizon control measurements become positive. This supplies the first natural-period witness outside `rule_109` and limits the earlier center-mediation necessity claim to the Fase 55 catalog; robustness remains partial and background-conditioned. Fases 89--90 then audit the historical detector cap globally. Replaying all 5,783,040 configurations preserves 445,897 unique historical positives and confirms 3,296 additional stationary oscillators with periods 18..120; none changes the `T=12` cohort used by the later causal audits. This converts the old `T<=16` negative label into an explicitly detector-bounded claim. Fases 91--112 then quotient those cases into 192 strict attractor classes, construct 48 complete Q8 intervention cubes, and derive the exact unit cut-coverage law on 43,425 exposures. The pair audit resolves minimum cardinality two in 69/126 vertex and 68/139 edge strata. Exhaustive audits of 3,061,466 triples and 24,362,850 quadruples resolve every remaining stratum: vertex minima partition as 69/41/16 and edge minima as 68/40/31 across cardinalities 2/3/4. The resulting 1,476 verified rescues form a 12-class Hamming-1 motif atlas and compile exactly into 265 sparse unit-cost QUBOs with no hardware or advantage claim. A conditioned audit then isolates `k2` as an empirically variable motif and, in 122 rescues across 24 mixed instances, proves the fixed partition `A_V+A_G+A_R=14`: external-attachment rescues allocate less connectivity toward `v_i` and more toward `G_min` on average. The result is descriptive of the frozen census and is not an externally validated classifier. Full results are in Sections 7.20-7.42.

Each extension is a controlled experiment with the same measurement protocol; only the IC or background definition changes.

10.3 PySR symbolic regression

The decision-tree calibration for `frontera_temporal` and `temporal_scale_stability` (Section 4) provides thresholds but not formulas. PySR symbolic regression on the fragility spectrum (`f_total`, `f_core`, `f_gap` as functions of rule properties and IC metrics) is now technically available, but the first atlas-wide run did not

produce a compact paper-worthy formula. The next useful symbolic-regression target would require mechanism features rather than aggregate atlas features alone.

10.4 Figures

The following six figures are planned for the preprint draft:

1. **World taxonomy table** -- the full 20-world atlas with categories, law coverage symbols, and fragility columns, formatted as a paper-ready table.
2. **Law coverage matrix** -- the `yes / partial / - / ?` matrix from `outputs/world_taxonomy/law_map.md`, rendered as a heatmap or binary grid.
3. **`f_total` / `f_core` spectrum** -- a two-axis scatter or bar chart showing all measured worlds positioned by `f_total` and `f_core`, with the four fragility mechanisms labeled.
4. **`rule_108` oscillator motif** -- a space-time diagram of the `## <-> ###` two-step cycle, showing several periods on a quiescent background.
5. **`rule_54` gate and observer non-equivariance** -- a dual figure: the Fase 13 noise-gate crossing diagram (reference dedup vs perturbed dedup) alongside the Fase 19 per-position dedup variation (15..24 across `k=0..63`).
6. **Moving oscillator space-time diagram** -- dual panel showing `rule_20` (right-moving) and `rule_6` (left-moving) over 8 time steps, with active cells dark and quiescent background light. Illustrates the `[0]` `<->` `[0,1]` glider cycle and the `+/-2` drift per period.

References

- [Wolfram2002] Stephen Wolfram. *A New Kind of Science*. Wolfram Media, Champaign, IL, 2002. <https://www.wolframscience.com/nks/>
- [Cook2004] Matthew Cook. "Universality in Elementary Cellular Automata." *Complex Systems* 15(1): 1-40, 2004. https://www.complex-systems.com/abstracts/v15_i01_a01/
- [Udrescu2020] Silviu-Marian Udrescu and Max Tegmark. "AI Feynman: A physics-inspired method for symbolic regression." *Science Advances* 6(16): eaay2631, 2020. <https://doi.org/10.1126/sciadv.aay2631>
- [Schmidt2009] Michael Schmidt and Hod Lipson. "Distilling Free-Form Natural Laws from Experimental Data." *Science* 324(5923): 81-85, 2009. <https://doi.org/10.1126/science.1165893>
- [Lu2024] Chris Lu, Cong Lu, Robert Tjarko Lange, Jakob Foerster, Jeff Clune, and David Ha. "The AI Scientist: Towards Fully Automated Open-Ended Scientific Discovery." arXiv:2408.06292, 2024. <https://arxiv.org/abs/2408.06292>
- [Crutchfield1989] James P. Crutchfield and Karl Young. "Inferring Statistical Complexity." *Physical Review Letters* 63(2):105-108, 1989. <https://doi.org/10.1103/PhysRevLett.63.105>
- [Shalizi2001] Cosma Rohilla Shalizi and James P. Crutchfield. "Computational Mechanics: Pattern and Prediction, Structure and Simplicity." *Journal of Statistical Physics* 104:817-879, 2001. <https://doi.org/10.1023/A:1010388907793>
- [Langton1990] Christopher G. Langton. "Computation at the Edge of Chaos: Phase Transitions and Emergent Computation." *Physica D: Nonlinear Phenomena* 42(1-3):12-37, 1990. [https://doi.org/10.1016/0167-2789\(90\)90064-V](https://doi.org/10.1016/0167-2789(90)90064-V)

- [Mitchell1993] Melanie Mitchell, Peter T. Hraber, and James P. Crutchfield. "Revisiting the Edge of Chaos: Evolving Cellular Automata to Perform Computations." *Complex Systems* 7:89-130, 1993. <https://content.wolfram.com/sites/13/2018/02/07-2-1.pdf>
- [Hanson1997] James E. Hanson and James P. Crutchfield. "Computational Mechanics of Cellular Automata: An Example." *Physica D: Nonlinear Phenomena* 103(1-4):169-189, 1997. [https://doi.org/10.1016/S0167-2789\(96\)00259-X](https://doi.org/10.1016/S0167-2789(96)00259-X)
- [Lindgren1988] Kristian Lindgren and Mats G. Nordahl. "Complexity Measures and Cellular Automata." *Complex Systems* 2(4), 1988. https://www.complex-systems.com/abstracts/v02_i04_a02/
- [Boccara1991] Nino Boccara, J. Nasser, and M. Roger. "Particlelike Structures and Their Interactions in Spatiotemporal Patterns Generated by One-Dimensional Deterministic Cellular-Automaton Rules." *Physical Review A* 44(2):866-875, 1991. <https://doi.org/10.1103/PhysRevA.44.866>
- [Martinez2006] Genaro Juarez Martinez, Andrew Adamatzky, and Harold V. McIntosh. "Phenomenology of Glider Collisions in Cellular Automaton Rule 54 and Associated Logical Gates." *Chaos, Solitons & Fractals* 28(1):100-111, 2006. <https://doi.org/10.1016/j.chaos.2005.05.013>
- [Redeker2010] Markus Redeker. "A Language for Particle Interactions in One-dimensional Cellular Automata." arXiv:1012.0158, 2010. <https://arxiv.org/abs/1012.0158>
- [Cranmer2023] Miles Cranmer. "Interpretable Machine Learning for Science with PySR and SymbolicRegression.jl." arXiv:2305.01582, 2023. <https://doi.org/10.48550/arXiv.2305.01582>
- [Elhage2021] Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Nova DasSarma, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Andy Jones, Jackson Kernion, Liane Lovitt, Kamal Ndousse, Dario Amodei, Tom Brown, Jack Clark, Jared Kaplan, Sam McCandlish, and Chris Olah. "A Mathematical Framework for Transformer Circuits." Transformer Circuits Thread, 2021. <https://transformer-circuits.pub/2021/framework/index.html>